DURHAM UNIVERSITY

DEPARTMENT OF COMPUTER SCIENCE

MASTER OF DATA SCIENCE

COMP42415
Text Mining and Language Analytics Workshops

Dr Stamos KATSIGIANNIS

2023-24

Preface

Welcome to the workshop manual for the *COMP42415 Text Mining and Language Analytics* course of Durham University. This manual has been meticulously crafted to serve as your comprehensive guide to the fascinating world of Natural Language Processing (NLP). As language plays a pivotal role in our daily lives, understanding and harnessing the power of natural language is increasingly crucial in various fields, ranging from artificial intelligence and data science to linguistics and computational linguistics.

This workshop is designed to provide you with a hands-on and in-depth exploration of NLP, equipping you with the knowledge and skills necessary to navigate the intricacies of natural language and to build practical applications. Whether you are a seasoned professional seeking to enhance your expertise or a beginner eager to delve into the realm of NLP, this manual is tailored to meet your needs.

Key Features of this Workshop Manual:

1. Comprehensive Coverage: This manual covers a wide array of topics, from the foundational concepts of NLP to advanced techniques and applications. Each section is structured to build upon the previous one, ensuring a logical progression of learning.

2. Hands-On Exercises: Learning by doing is at the core of this workshop. Throughout the manual, you will find hands-on exercises that allow you to apply the concepts you learn in real-world scenarios. These exercises are designed to reinforce your understanding and enhance your practical skills.

3. Real-World Applications: NLP has diverse applications, and this manual provides insights into how NLP techniques are used in real-world scenarios.

4. Do-It-Yourself: This manual is designed to support self-learning, providing detailed examples and explanations for the presented NLP techniques.

By the end of this workshop, you will have gained a solid understanding of NLP concepts, developed practical skills in implementing NLP solutions, and be well-prepared to tackle challenges in the ever-evolving landscape of NLP.

We hope you find this workshop manual informative, engaging, and instrumental in your quest to master Natural Language Processing.

Best wishes for a rewarding learning experience!

Dr Stamos Katsigiannis

Required Python packages

Required python version: This workshop manual has been tested on <u>Python 3.11</u>.

The following Python packages are required:

- re
- nltk
- numpy
- scipy
- math
- scikit-learn
- matplotlib
- seaborn
- pandas
- tensorflow
- pickle

Contents

Workshop 1: Text query with Regular Expressions

1.1 Regular Expressions Definition

A regular expression (shortened as regex or regexp) is a sequence of characters that define a search pattern. Usually such patterns are used by string-searching algorithms for "find" or "find and replace" operations on strings, or for input validation.

1.2 The "re" Python package

Python has a built-in package called "re", which can be used to work with Regular Expressions. Let's import this package and use a regular expression to check whether the sentence "*I started studying this year at Durham University*" ends with the word "University" or with the word "school".

1.2.1 Example

```python
import re # Import the re package

txt = "I started studying this year at Durham University"

x1 = re.search("University$", txt) # Returns a Match object if there is a match anywhere in the
    string with the regex
x2 = re.search("school$", txt)

print("x1:",x1)
print("x2:",x2)

if(x1):
    print("The text ends with 'University'")
else:
    print("The text does not end with 'University'")

if(x2):
    print("The text ends with 'school'")
else:
    print("The text does not end with 'school'")
```

The output will look like:

```
x1: <_sre.SRE_Match object; span=(39, 49), match='University'>
x2: None
The text ends with 'University'
The text does not end with 'school'
```

1.2.2 RegEx functions in "re"

The re module offers a set of functions that allows us to search a string for a match:

Function	Description
findall(*args*)	Returns a list containing all matches
search(*args*)	Returns a Match object if there is a match anywhere in the string
match(*args*)	Returns a Match object if there is a match starting at the beginning of the string
split(*args*)	Returns a list where the string has been split at each match
sub(*args*)	Replaces one or many matches with a string

1.2.3 Regular expression metacharacters

As you can see in our first example, we used the character "$" in order to indicate that a text matching the regular expression should end with the string preceding the character "$". For example, the regular expression "car$" indicates that the text should and with the string "car". In this case, "$" is considered as a metacharacter, i.e. a character with a special meaning. Below are the metacharacters supported by the "re" package:

Character	Description	Example
[ ]	A set of characters	"[a-f]"
\	Signals a special sequence (can also be used to escape special characters)	"\s"
.	Any character (except newline character)	"uni..rsity"
^	Starts with	"^She"
$	Ends with	"John$"
*	Zero or more occurrences	"o*"
+	One or more occurrences	"l+"
?	Matches 0 or 1 repetitions of the preceding regex	ab? will match either "a" or "ab"
{}	Exactly the specified number of occurrences	"o{2}"
|	Either or	"he|she"
()	Capture and group	

The first metacharacters we'll look at are [and]. They're used for specifying a character class, which is a set of characters that you wish to match. Characters can be listed individually, or a range of characters can be indicated by giving two characters and separating them by a "-". For example, [abc] will match any of the characters a, b, or c; this is the same as [a-c], which uses a range to express the same set of characters. If you wanted to match only lowercase letters, your regex would be [a-z].

Metacharacters are not active inside classes. For example, [akm$] will match any of the characters "a", "k", "m", or "$". "$" is usually a metacharacter, but inside a character class it is stripped of its special nature.

You can match the characters not listed within the class by complementing the set. This is indicated by including a "^" as the first character of the class. For example, [^5] will match any character except "5". If the caret appears elsewhere in a character class, it does not have special meaning. For example: [5^] will match either a "5" or a "^".

Perhaps the most important metacharacter is the backslash, \. As in Python string literals, the backslash can be followed by various characters to signal various special sequences. It is also used to escape all the metacharacters so you can still match them in patterns. For example, if you need to match a [or \, you can precede them with a backslash to remove their special meaning: \[or \\.

Some of the special sequences beginning with "\" represent predefined sets of characters that are often useful, such as the set of digits, the set of letters, or the set of anything that isn't a white space.

1.2.4 Regular expression special sequences

Let's see some of the main regular expression special sequences. For a more detailed list, please refer to https://docs.python.org/3/library/re.html#re-syntax.

These sequences can be included inside a character class. For example, [\s,.] is a character class that will match any white space character, or "," or ".".

Sequence	Description
\d	Matches any decimal digit. Equivalent to the class [0-9].
\D	Matches any non-digit character. Equivalent to the class [^0-9].
\s	Matches any white space character. Equivalent to the class [\t\n\r\f\v].
\S	Matches any non-white space character. Equivalent to the class [^\t\n\r\f\v].
\w	Matches any alphanumeric character. Equivalent to the class [a-zA-Z0-9_].
\W	Matches any non-alphanumeric character. Equivalent to the class [^a-zA-Z0-9_].

1.2.5 Repeating regular expressions

Being able to match varying sets of characters is the first thing regular expressions can do that isn't already possible with the methods available on strings. However, if that was the only additional capability of regexes, they wouldn't be much of an advance. Another capability is that you can specify that portions of the regular expression must be repeated a certain number of times.

The first metacharacter for repeating things that we'll look at is "*". "*" does not match the literal character "*", but it specifies that the previous character can be matched zero or more times, instead of exactly once.

Example: "do*g" will match "dg" (zero "o" characters), "dog" (one "o" character), "doooog" (four "o" characters), and so on.

Repetitions such as "*","+", and "?" are greedy. When repeating a regular expression, the matching engine will try to repeat it as many times as possible. If later portions of the pattern don't match, the matching engine will then back up and try again with fewer repetitions. If this behaviour is undesirable, you can add "?" after the qualifier ("*?","+?", "??") to make it perform the match in non-greedy or minimal fashion, i.e. as few characters as possible will be matched.

Step-by-step example
Let's consider the expression a[bcd]*b. This matches the letter "a", zero or more letters from the class [bcd], and finally ends with a "b". Now imagine matching this regular expression against the string "abcbd".

Step	Matched	Explanation
1	a	The "a" in the regex matches.
2	abcbd	The engine matches "[bcd]*", going as far as it can, which is to the end of the string.
3	**FAILED**	The engine tries to match "b", but the current position is at the end of the string, so it fails.
4	abcb	Back up, so that "[bcd]*" matches one less character.
5	**FAILED**	Try "b" again, but the current position is at the last character, which is a "d".
6	abc	Back up again, so that "[bcd]*" is only matching "bc".
7	abcb	Try "b" again. This time the character at the current position is "b", so it succeeds.

The end of the regular expression has now been reached, and it has matched "abcb". This demonstrates how the matching engine goes as far as it can at first, and if no match is found it will then progressively back up and retry the rest of the regular expression again and again. It will back up until it has tried zero matches for "[bcd]*", and if that subsequently fails, the engine will conclude that the string does not match the regex at all.

Another repeating metacharacter is "+", which matches one or more times. Pay careful attention to the difference between "*" and "+". "*" matches zero or more times, so whatever's being repeated may not be present at all, while "+" requires at least one occurrence.

Example: "do+g" will match "dog" (one "o" character), "dooog" (three "o" characters), and so on, but will not match "dg" (zero "o" characters).

There are two more repeating qualifiers. The question mark character "?" matches either once or zero times. Think of it as marking something as being optional.

Example: "pre-?processing" matches either "preprocessing" or "pre-processing".

The most complicated repeated qualifier is "{m,n}", where m and n are decimal integers. This qualifier means there must be at least m repetitions, and at most n. For example, "a/{1,3}b" will match "a/b", "a//b", and "a///b", but it will not match "ab", which has no slashes, or "a////b", which has four slashes. You can omit either m or n. In this case, default values for m or n are used. Omitting m is interpreted as a lower limit of 0, while omitting n results in an upper bound of infinity.

Note: Some qualifiers are interchangeable. For example "{0,}" is the same as "*", "{1,}" is the same as "+", and "{0,1}" is the same as "?". "*", "+", and "?" make the regular expression easier to read, so try to use them if possible.

1.3   Using Regular Expressions to match strings

1.3.1   Matching example

Let's use the text "*I started studying this year at Durham University*" again and find out whether the string "at" or the string "in" exists in the text.

```python
txt = "I started studying this year at Durham University"

x = re.search("at|in", txt)
print(x.string) # Returns the string passed into the function
print(x.span()) # Returns a tuple containing the start, and end positions of the match
print(x.group()) # Returns the part of the string where there was a match

print(txt[x.span()[0]:x.span()[1]]) # Print the content of the string at the positions of the match
```

The output will look like:

```
I started studying this year at Durham University
(15, 17)
in
in
```

As you can see, there was a match to our regex at the character with index 15 (counting starts from 0), ending at the character with index 17. Indeed, the string "in" was found within the word "studying".

However, if you read the input text, there should have been a second match for the word "at" but only the first match was returned. Note that if there is more than one match, only the first occurrence of the match will be returned by the search() function! We can use the findall() function to get a list of all matches in the order they are found.

```python
x = re.findall("at|in", txt)
for match in x:
    print(match)
```

The output will look like:

```
in
at
```

As expected, the findall() function returned two matches, "in" and "at".

Consider the string "*stp stop stoop stooooop stooooooooooop*". Let's find matches where the character "o" appears one or more times and matches where the character "o" appears two times only.

```python
x = re.findall("o+", "stp stop stoop stooooop stooooooooooop")
print("One or more occurrences of 'o':")
for match in x:
    print(match)

x = re.findall("o{2}", "stp stop stoop stooooop stooooooooooop")
print("\nTwo occurrences of 'o':")
for match in x:
    print(match)
```

The output will look like:

```
One or more occurrences of 'o':
o
oo
oooooo
oooooooooo

Two occurrences of 'o':
oo
oo
oo
oo
oo
oo
oo
oo
oo
```

Note that when looking for two occurrences of "o", the findall() function returned 9 matches that correspond to the 9 non-overlapping matches from the start (left) to the end (right) of the input text: "*stp stop st(oo)p st(oo)(oo)(oo)p st(oo)(oo)(oo)(oo)(oo)op*"

1.3.2 Matching to validate strings

Let's use regular expressions to check the validity of various strings. Consider an identification number that should consist of exactly 10 digits (from 0 to 9), and the strings: "0123456789", "12345", "0000a00005", "+000001111", "00000011115", "2030405060". How can we check whether these strings are valid identification numbers?

```python
text = list()
text.append("0123456789")
text.append("12345")
text.append("0000a00005")
text.append("+000001111")
text.append("00000011115")
text.append("2030405060")

regex = "[0-9]{10}"

result = list()
for t in text:
    x = re.match(regex, t)
    if(x != None):
        print(t,"->",x.group())
    else:
        print(t,"-> No match")
    result.append(x)

for i in range(len(text)):
    if(result[i]!=None and result[i].group()==text[i]):
        print(text[i],"is a valid identification number")
    else:
        print(text[i],"is NOT a valid identification number")
```

The output will look like:

```
0123456789 -> 0123456789
12345 -> No match
0000a00005 -> No match
+000001111 -> No match
00000011115 -> 0000001111
2030405060 -> 2030405060
0123456789 is a valid identification number
12345 is NOT a valid identification number
0000a00005 is NOT a valid identification number
```

```
+000001111 is NOT a valid identification number
00000011115 is NOT a valid identification number
2030405060 is a valid identification number
```

Notice that string "00000011115" consists of 11 numerical digits, thus the regular expression matches the subset "0000001111". However, this is not a valid identification number according to the specification above. When validating input, remember to check whether the matched string is equal to the query string.

1.3.3   Credit card number validation

Let's try to validate whether the following strings are valid VISA or Mastercard credit card numbers: "1000000000000", "4000000000000", "5000000000000", "50000000000000000", "50000a0000000c000", "40123456789". VISA credit card numbers should start with a 4 and have 13 or 16 digits. Mastercard credit card numbers start with a 5 and have 16 digits.

```python
text = list()
text.append("1000000000000") # 13 digits - Not valid
text.append("4000000000000") # 13 digits - Valid VISA
text.append("5000000000000") # 13 digits - Not valid
text.append("5000000000000000") # 16 digits - Valid Mastercard
text.append("50000a0000000c000") # Not valid, contains letters
text.append("40123456789") # 11 digits - Not valid

regex = "(5[0-9]{15})|(4([0-9]{12}|[0-9]{15}))" # Number 5 followed by 15 digits OR number 4 followed
    by either 12 or 15 digits

result = list()
for t in text:
    x = re.match(regex, t)
    if(x != None):
        print(t,"->",x.group())
    else:
        print(t,"-> No match")
    result.append(x)

print("")
for i in range(len(text)):
    if(result[i]!=None and result[i].group()==text[i]):
        print(text[i],"is a valid VISA or Mastercard number")
    else:
        print(text[i],"is NOT a valid VISA or Mastercard number")
```

The output will look like:

```
1000000000000 -> No match
4000000000000 -> 4000000000000
5000000000000 -> No match
5000000000000000 -> 5000000000000000
50000a0000000c000 -> No match
40123456789 -> No match

1000000000000 is NOT a valid VISA or Mastercard number
4000000000000 is a valid VISA or Mastercard number
5000000000000 is NOT a valid VISA or Mastercard number
5000000000000000 is a valid VISA or Mastercard number
50000a0000000c000 is NOT a valid VISA or Mastercard number
40123456789 is NOT a valid VISA or Mastercard number
```

Let's analyse the regex that we used. Both VISA and Mastercard numbers start with a specific digit but Mastercard has exactly 16 digits in total, while VISA can have either 13 or 16 digits. Let's first create a regex for each case separately. For Mastercard, it should be "5[0-9]{15}", the digit 5 followed by exactly 15 digits (0-9), for a total of 16 digits. For VISA, it should be "4([0-9]{12}|[0-9]{15})", the digit 4 followed by either exactly 12 digits, for a total of 13 digits, or exactly 15 digits, for a total of 16 digits. Then, to include both the VISA and the Mastercard cases in our final regex, we can enclose each regex within parentheses and combine them with the OR ("|") operator.

Note that the expression [0-9] could be switched to [\d]:

```python
regex = "(5[\d]{15})|(4([\d]{12}|[\d]{15}))" # Number 5 followed by 15 digits OR number 4 followed by
    either 12 or 15 digits

result = list()
for t in text:
    x = re.match(regex, t)
    if(x != None):
        print(t,"->",x.group())
    else:
        print(t,"-> No match")
    result.append(x)

print("")
for i in range(len(text)):
    if(result[i]!=None and result[i].group()==text[i]):
        print(text[i],"is a valid VISA or Mastercard number")
    else:
        print(text[i],"is NOT a valid VISA or Mastercard number")
```

The output will look like:

```
1000000000000 -> No match
4000000000000 -> 4000000000000
5000000000000 -> No match
5000000000000000 -> 5000000000000000
50000a0000000c000 -> No match
40123456789 -> No match

1000000000000 is NOT a valid VISA or Mastercard number
4000000000000 is a valid VISA or Mastercard number
5000000000000 is NOT a valid VISA or Mastercard number
5000000000000000 is a valid VISA or Mastercard number
50000a0000000c000 is NOT a valid VISA or Mastercard number
40123456789 is NOT a valid VISA or Mastercard number
```

1.3.4   Validation of United Kingdom's National Insurance numbers (NINO)

According to the rules for validating UK national insurance numbers[1], a NINO is made up of 2 letters, 6 numbers and a final letter, which is always A, B, C, or D. It looks something like this: QQ 12 34 56 A. The characters D, F, I, Q, U, and V are not used as either the first or second letter of a NINO prefix. The letter O is not used as the second letter of a prefix.

Let's create the required regex step-by-step and validate the following strings: "AA 123456 B", "AO 123456 B", "AQ123456 B", "QQ 123456 B", "AA 123456 X", "AA 12 34 56 B", "AA123456B", "AA12345B", "A 123456 B", "A 123456 B". We must also take white spaces into consideration. Let's consider the following two ways of writing a NINO: AA123456A and AA 12345 A.

1. The first letter should be any of A, B, C, E, G, H, J, K, L, M, N, O, P, R, S, T, W, X, Y:
   (A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)

2. The second letter should be any of A, B, C, E, G, H, J, K, L, M, N, P, R, S, T, W, X, Y:
   (A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)**(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)**

3. The third letter can optionally be a white space character:
   (A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)**[\s]?**

4. Then, exactly 6 digits (0-9) are required:
   (A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)[\s]?**[0-9]{6}**

5. The next letter can optionally be a white space character:
   (A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)[\s]?[0-9]{6}**[\s]?**

---

[1]https://www.gov.uk/hmrc-internal-manuals/national-insurance-manual/nim39110

6. The final letter must be one of A, B, C, or D:

(A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)[\s]?[0-9]{6}[\s]?(**A**|**B**|**C**|**D**)

```python
text = list()
text.append("AA 123456 B") # Valid
text.append("AO 123456 B") # Not valid
text.append("AQ123456 B") # Not valid
text.append("QQ 123456 B") # Not valid
text.append("AA 123456 X") # Not valid
text.append("AA 12 34 56 B") # Not valid
text.append("AA123456B") # Valid
text.append("AA12345B") # Not valid
text.append("A 123456 B") # Not valid
text.append("A 123456 B") # Not valid

regex =
    "(A|B|C|E|G|H|J|K|L|M|N|O|P|R|S|T|W|X|Y)(A|B|C|E|G|H|J|K|L|M|N|P|R|S|T|W|X|Y)[\s]?[0-9]{6}[\s]?(A|B|C|D)"

result = list()
for t in text:
    x = re.match(regex, t)
    if(x != None):
        print(t,"->",x.group())
    else:
        print(t,"-> No match")
    result.append(x)

print("")
for i in range(len(text)):
    if(result[i]!=None and result[i].group()==text[i]):
        print("VALID\t",text[i])
    else:
        print("-----\t",text[i])
```

The output will look like:

```
AA 123456 B -> AA 123456 B
AO 123456 B -> No match
AQ123456 B -> No match
QQ 123456 B -> No match
AA 123456 X -> No match
AA 12 34 56 B -> No match
AA123456B -> AA123456B
AA12345B -> No match
A 123456 B -> No match
A  123456 B -> No match

VALID AA 123456 B
----- AO 123456 B
----- AQ123456 B
----- QQ 123456 B
----- AA 123456 X
----- AA 12 34 56 B
VALID AA123456B
----- AA12345B
----- A 123456 B
----- A  123456 B
```

1.3.5   Validation of hexadecimal numbers

Let's use regular expressions to check whether a string corresponds to a hexadecimal number. Consider the strings "xAF1400BD", "1299ab32", "xFF00FF5R", "0xaa00bb". How can we check if these strings are representations of hexadecimal numbers? Remember that valid hexadecimal digits are [0,1,2,3,4,5,6,7,8,9,A,B,C,D,E,F, a,b,c,d,e,f] and that in computers, hexadecimal numbers may be denoted with an "x" or "0x" (either lowercase

or uppercase) in the beginning. For example, the hexadecimal number FFF, can also be written as fff, 0xfff, 0XFFF, xfff, XFFF and can also have mixed lowercase and uppercase letters.

```python
text = list()
text.append("xAF1400BD") # Valid
text.append("1299ab32") # Valid
text.append("xFF00FF5R") # Not valid - Character R is not a valid headecimal digit
text.append("0xaa00bb") # Valid
text.append("xaa00bb4657AB000922334bce111A") # Valid
text.append("0xfff") # Valid
text.append("xFfF") # Valid
text.append("AAA") # Valid
text.append("ALA") # Not valid - Character L is not a valid headecimal digit

regex = "(0x|0X|x|X)?[0-9a-fA-F]+" # One optional occurence of 0x, 0X, x, or X, followed by at least
    one digit from 0 to 9, or lowercase letter from a to f, or uppercase letter from A to F

result = list()
for t in text:
    x = re.match(regex, t)
    if(x != None):
        print(t,"->",x.group())
    else:
        print(t,"-> No match")
    result.append(x)

print("")
for i in range(len(text)):
    if(result[i]!=None and result[i].group()==text[i]):
        print("VALID HEX\t",text[i])
    else:
        print("---------\t",text[i])
```

The output will look like:

```
xAF1400BD -> xAF1400BD
1299ab32 -> 1299ab32
xFF00FF5R -> xFF00FF5
0xaa00bb -> 0xaa00bb
xaa00bb4657AB000922334bce111A -> xaa00bb4657AB000922334bce111A
0xfff -> 0xfff
xFfF -> xFfF
AAA -> AAA
ALA -> A

VALID HEX    xAF1400BD
VALID HEX    1299ab32
---------    xFF00FF5R
VALID HEX    0xaa00bb
VALID HEX    xaa00bb4657AB000922334bce111A
VALID HEX    0xfff
VALID HEX    xFfF
VALID HEX    AAA
---------    ALA
```

Note that in the case of the "ALA" string, the regex found the match "A", which is a valid hexadecimal number, but the full string "ALA" is not a valid hexadecimal number. When validating input, remember to check that the matched string is equal to the query string.

1.4 Using Regular Expressions to search elements in files

1.4.1 Parsing XML files

Let's use regular expressions to parse an XML file. We will use the movies.xml file which contains the titles and release dates of 5 movies. We will use a regular expression to retrieve all the movie titles from the file. First,

copy the file "movies.xml" to your current working directory or retrieve the absolute path of the file. Then open the file, load its contents into a variable and close the file.

```python
f = open("movies.xml", "r") # Opens the file for reading only ("r")
text = f.read() # Store the contents of the file in variable "text". read() returns all the contents
    of the file
f.close() # Close the file
print(text) # Print the contents of variable "text"
```

The output will look like:

```xml
<movies>
  <movie id="1">
    <title>And Now for Something Completely Different</title>
    <released>1971</released>
  </movie>
  <movie id="2">
    <title>Monty Python and the Holy Grail</title>
    <released>1974</released>
  </movie>
  <movie id="3">
    <title>Monty Python's Life of Brian</title>
    <released>1979</released>
  </movie>
  <movie id="4">
    <title>Monty Python Live at the Hollywood Bowl</title>
    <released>1982</released>
  </movie>
  <movie id="5">
    <title>Monty Python's The Meaning of Life</title>
    <released>1983</released>
  </movie>
</movies>
```

As you can see above, all movie titles in the movies.xml XML file are enclosed within the <title> and </title> tags. Let's use a regular expression to match all the strings that are enclosed within these tags.

```python
regex = "<title>.*</title>" # <title>, followed by any number of characters except for the newline
    character, followed by </title>

x = re.findall(regex, text) # Find all matches of the regex

print(x,"\n")

print("Movie titles from movies.xml:")
for title in x:
    title = title.replace("<title>","") # Remove <title> by replacing it with empty string
    title = title.replace("</title>","") # Remove </title> by replacing it with empty string
    print(title)
```

The output will look like:

```
['<title>And Now for Something Completely Different</title>', '<title>Monty Python and the Holy
    Grail</title>', "<title>Monty Python's Life of Brian</title>", '<title>Monty Python Live at the
    Hollywood Bowl</title>', "<title>Monty Python's The Meaning of Life</title>"]

Movie titles from movies.xml:
And Now for Something Completely Different
Monty Python and the Holy Grail
Monty Python's Life of Brian
Monty Python Live at the Hollywood Bowl
Monty Python's The Meaning of Life
```

1.4.2 Parsing HTML files

Let's read the links.html HTML file and use regular expressions to find all links within the file. Remember that in HTML, links are denoted using the <a> and tags and the link url is provided using the "href" attribute within the <a> tag. For example a link for the main website of Durham University would be:

```
<a href="https://www.dur.ac.uk/">Durham University</a>
```

```python
f = open("links.html", "r") # Opens the file for reading only ("r")
text = f.read() # Store the contents of the file in variable "text". read() returns all the contents
    of the file
f.close() # Close the file
print(text,"\n") # Print the contents of variable "text"

regex = '<a href=".*"'

x = re.findall(regex, text) # Find all matches of the regex

print(x,"\n")

print("Links from links.xml:")
for link in x:
    link = link.replace('<a href="',"") # Remove <a href=" by replacing it with empty string
    link = link.replace('"',"") # Remove " by replacing it with empty string
    print(link)
```

The output will look like:

```
<!doctype html>
<html lang="en-GB">

<head>
   <meta name="viewport" content="width=device-width, initial-scale=1">
   <title>Test page for Text Mining and Language Analytics</title>
</head>

<body>
   <h1>Test page for Text Mining and Language Analytics</h1>
   <a href="https://www.durham.ac.uk/departments/academic/computer-science/">Link 1</a><br/>
   <a href="https://www.gov.uk/government/organisations/hm-revenue-customs">Link 2</a><br/>
   <a href="https://www.gov.uk/">Link 3</a><br/>
   <a href="https://www.dur.ac.uk/">Link 4</a><br/>
   <a href="https://www.nhs.uk/">Link 5</a><br/>
</body>

</html>

['<a href="https://www.durham.ac.uk/departments/academic/computer-science/"', '<a
    href="https://www.gov.uk/government/organisations/hm-revenue-customs"', '<a
    href="https://www.gov.uk/"', '<a href="https://www.dur.ac.uk/"', '<a href="https://www.nhs.uk/"']

Links from links.xml:
https://www.durham.ac.uk/departments/academic/computer-science/
https://www.gov.uk/government/organisations/hm-revenue-customs
https://www.gov.uk/
https://www.dur.ac.uk/
https://www.nhs.uk/
```

Note than we used single quotes to denote strings that included a double quote character.

As you can see above, we successfully retrieved the urls from links.html. Nevertheless, please note that using regular expressions is not the best approach for parsing HTML files due to the flexibility of HTML syntax. Solutions like XPath (`https://developer.mozilla.org/en-US/docs/Web/API/XPathExpression`) are more suitable for HTML parsing.

1.4.3   Parsing raw text

The file emails.txt contains a list of emails from various domains. Let's use regular expressions to find the emails from the durham.ac.uk domain.

```python
f = open("emails.txt", "r") # Opens the file for reading only ("r")
text = f.read() # Store the contents of the file in variable "text". read() returns all the contents
    of the file
f.close() # Close the file
print(text,"\n") # Print the contents of variable "text"

regex = "[0-9a-zA-Z!#$%&'*+-/=?^_`{|}~.]+@durham.ac.uk"

x = re.findall(regex, text) # Find all matches of the regex

print(x,"\n")

print("Emails from durham.ac.uk:")
for email in x:
    print(email)
```

The output will look like:

```
john+acme.co@hotmail.com
bob@gmail.com
tom@durham.ac.uk
jerry@durham.ac.uk
scrooge@durham.ac.uk
donald@yahoo.co.uk
huey@yahoo.co.uk
dewey@gmail.com
louie.duck@durham.ac.uk
gyro.gearloose@yahoo.co.uk
bart@yahoo.co.uk
homer@gmail.com
stan@hotmail.com
kyle-broflovski@durham.ac.uk
eric@yahoo.co.uk
kenny@gmail.com
butters@durham.ac.uk
wendy@hotmail.com
randy_marsh@durham.ac.uk
chef@gmail.com

['tom@durham.ac.uk', 'jerry@durham.ac.uk', 'scrooge@durham.ac.uk', 'louie.duck@durham.ac.uk',
    'kyle-broflovski@durham.ac.uk', 'butters@durham.ac.uk', 'randy_marsh@durham.ac.uk']

Emails from durham.ac.uk:
tom@durham.ac.uk
jerry@durham.ac.uk
scrooge@durham.ac.uk
louie.duck@durham.ac.uk
kyle-broflovski@durham.ac.uk
butters@durham.ac.uk
randy_marsh@durham.ac.uk
```

As you can see above, we retrieved all emails from the durham.ac.uk email.

1.5   Using Regular Expressions to substitute strings

1.5.1   Substitution example

Let's now convert any string that matches the "at|in" regex in the text "I started studying this year at Durham University" with the string "FOO". To achieve this, we are going to use the sub() function.

```
txt = "I started studying this year at Durham University"
x = re.sub("at|in","FOO", txt)
print(x)
```

The output will look like:

```
I started studyFOOg this year FOO Durham University
```

As expected, two matches of the regex were converted to "FOO". What if we wanted only the first match to be substituted with "FOO"? We can add an additional argument in the sub() function indicating the number of substitutions we would like to make.

```
x = re.sub("at|in","FOO", txt,1)
print(x)
```

The output will look like:

```
I started studyFOOg this year at Durham University
```

1.5.2   Email domain substitution

Let's load again emails.txt and change the domain to "new.ac.uk" for all emails in the "durham.ac.uk", "gmail.com", and "yahoo.co.uk" domains.

```
f = open("emails.txt", "r") # Opens the file for reading only ("r")
text = f.read() # Store the contents of the file in variable "text". read() returns all the contents
    of the file
f.close() # Close the file
print("OLD EMAILS:")
print(text,"\n") # Print the contents of variable "text"

regex = "@durham.ac.uk|@gmail.com|@yahoo.co.uk"

print("NEW EMAILS:")
x = re.sub(regex,"@new.ac.uk", text)

print(x)
```

The output will look like:

```
OLD EMAILS:
john+acme.co@hotmail.com
bob@gmail.com
tom@durham.ac.uk
jerry@durham.ac.uk
scrooge@durham.ac.uk
donald@yahoo.co.uk
huey@yahoo.co.uk
dewey@gmail.com
louie.duck@durham.ac.uk
gyro.gearloose@yahoo.co.uk
bart@yahoo.co.uk
homer@gmail.com
stan@hotmail.com
kyle-broflovski@durham.ac.uk
eric@yahoo.co.uk
kenny@gmail.com
butters@durham.ac.uk
wendy@hotmail.com
randy_marsh@durham.ac.uk
chef@gmail.com

NEW EMAILS:
john+acme.co@hotmail.com
```

```
bob@new.ac.uk
tom@new.ac.uk
jerry@new.ac.uk
scrooge@new.ac.uk
donald@new.ac.uk
huey@new.ac.uk
dewey@new.ac.uk
louie.duck@new.ac.uk
gyro.gearloose@new.ac.uk
bart@new.ac.uk
homer@new.ac.uk
stan@hotmail.com
kyle-broflovski@new.ac.uk
eric@new.ac.uk
kenny@new.ac.uk
butters@new.ac.uk
wendy@hotmail.com
randy_marsh@new.ac.uk
chef@new.ac.uk
```

1.6   Exercises

Exercise 1.1 Validate the following identification numbers: "500011110000" (valid), "5000 1111 0000" (valid), "500001110000" (valid), "5000 0111 0000" (valid), "500021110000" (not valid), "5000 2111 0000" (not valid), "400001110000" (not valid), "4000 0111 0000" (not valid), "500011110009" (not valid), "5000 1111 0009" (not valid). A valid number should be 12 digits long, the first digit should always be 5, the 5th digit should be 0 or 1, and the last digit cannot be 8 or 9. The numbers can also be grouped in groups of four digits with a white space between groups.

Exercise 1.2 Redo the activity from Section 1.3.4 in order to also support the following writing of NINO strings: "AA 12 34 56 A"

Exercise 1.3 Use regular expressions to print a list of all the album titles, a list of all the artists, and the average price of all CDs in the cds.xml file.

Exercise 1.4 Create a regular expression that matches all the strings in the first column but none of those in the second column

affgfking	fgok
rafgkahe	a fgk
bafghk	affgm
baffgkit	afffhk
affgfking	fgok
rafgkahe	afg.K
bafghk	aff gm
baffgkit	afffhgk

Exercise 1.5 Use a regular expression to substitute the quantities of the bought items with "XX" in the following sentence: "*Yesterday we bought 120 packs of A4 paper, 5 bottles of ink, 10 boxes of paperclips, 200 notebooks, and 5.35 litres of fuel. The total cost for order #1290 was £1000.43.*". Note that the order number and cost must not be substituted. The resulting sentence must be: "*Yesterday we bought XX packs of A4 paper, XX bottles of ink, XX boxes of paperclips, XX notebooks, and XX litres of fuel. The total cost for order #1290 was £1000.43.*"

Workshop 2: Text pre-processing

2.1 NLTK corpora

The NLTK python package comes pre-loaded with a lot of corpora that can be used to experiment with text pre-processing.

2.1.1 Import NLTK

Let's load one of the available corpora in NLTK, called movie_reviews. The movie_reviews corpus contains documents consisting of reviews about movies and annotated in relation to their sentiment.

```python
import nltk # Import the NLTK library
nltk.download('movie_reviews') # Download movie_reviews from NLTK
from nltk.corpus import movie_reviews # Import the movie_reviews corpus from NLTK
```

2.1.2 Corpus file IDs

Let's check the file IDs in the movie_reviews corpus:

```python
movie_reviews.fileids() # List file-ids in the corpus
```

The output will look like:

```
['neg/cv000_29416.txt', 'neg/cv001_19502.txt', 'neg/cv002_17424.txt', 'neg/cv003_12683.txt',
    'neg/cv004_12641.txt', ...]
```

2.1.3 Corpus file categories

Let's examine what classes (categories) are included in the corpus:

```python
movie_reviews.categories() # List categories in the corpus
```

The output will look like:

```
['neg', 'pos']
```

This means that the movie_reviews corpus contains documents that have been annotated as belonging to two distinct classes, more specifically the Negative class (neg) and the Positive class (pos), in relation to the sentiment of the respective movie review.

2.1.4 Corpus words

Let's see the list of words in the movie_review corpus and print their number:

```python
print(movie_reviews.words())
length = len(movie_reviews.words())
print("Number of words in corpus: ", length)
```

The output will look like:

```
['plot', ':', 'two', 'teen', 'couples', 'go', 'to', ...]
Number of words in corpus: 1583820
```

2.1.5   Selecting files from specific category

Use the following in order to view the file IDs for files belonging to a specific category. Note that you can provide a list with categories, e.g ['neg','pos','other']

```
movie_reviews.fileids(['neg']) # List file ids with 'neg' category
movie_reviews.fileids(['pos']) # List file ids with 'pos' category
```

The output will look like:

```
['neg/cv000_29416.txt', 'neg/cv001_19502.txt', 'neg/cv002_17424.txt', 'neg/cv003_12683.txt',
'neg/cv004_12641.txt', 'neg/cv005_29357.txt', 'neg/cv006_17022.txt', 'neg/cv007_4992.txt', ...]
['pos/cv000_29590.txt', 'pos/cv001_18431.txt', 'pos/cv002_15918.txt', 'pos/cv003_11664.txt',
'pos/cv004_11636.txt', 'pos/cv005_29443.txt', 'pos/cv006_15448.txt', 'pos/cv007_4968.txt', ...]
```

2.1.6   Corpus sentences

To access the tokenised sentences included in the corpus:

```
nltk.download('punkt') # Download the Punkt sentence tokenizer from NLTK

movie_reviews.sents()
```

The output will look like:

```
[['plot', ':', 'two', 'teen', 'couples', 'go', 'to', 'a', 'church', 'party', ',', 'drink', 'and',
    'then', 'drive', '.'], ['they', 'get', 'into', 'an', 'accident', '.'], ...]
```

2.1.7   Accessing sentences from specific file

We can use the "fileids" argument in order to access the sentences from a specific file in the corpus:

```
movie_reviews.sents(fileids='pos/cv004_11636.txt')
```

The output will look like:

```
[['moviemaking', 'is', 'a', 'lot', 'like', 'being', 'the', 'general', 'manager', 'of', 'an', 'nfl',
    'team', 'in', 'the', 'post', '-', 'salary', 'cap', 'era', '--', 'you', '"', 've', 'got', 'to',
    'know', 'how', 'to', 'allocate', 'your', 'resources', '.'], ['every', 'dollar', 'spent', 'on',
    'a', 'free', '-', 'agent', 'defensive', 'tackle', 'is', 'one', 'less', 'dollar', 'than', 'you',
    'can', 'spend', 'on', 'linebackers', 'or', 'safeties', 'or', 'centers', '.'], ...]
```

2.1.8   Number of documents in each category

Let's see how many documents (files) does the movie_reviews corpus contain, and then how many documents of "neg" category and how many of "pos" category it contains:

```
documents = len(movie_reviews.fileids())
documents_neg = len(movie_reviews.fileids(['neg']))
documents_pos = len(movie_reviews.fileids(['pos']))
print("Number of documents: ", documents)
print("Number of documents in neg category: ", documents_neg)
print("Number of documents in pos category: ", documents_pos)
```

The output will look like:

```
Number of documents: 2000
```

```
Number of documents in neg category: 1000
Number of documents in pos category: 1000
```

As you can see here, the corpus is perfectly balanced between negative and positive sentiment reviews, having 1000 documents in each category.

2.1.9   Corpus raw text

To access the raw text of a specific file, use the ".raw()" function:

```python
rawtext = movie_reviews.raw('neg/cv002_17424.txt').strip()[:500] # strip() removes blank spaces in
    the beginning and end of a string. [:500] is used to only retrieve the first 500 characters of
    the string
print(rawtext)
```

The output will look like:

```
it is movies like these that make a jaded movie viewer thankful for the invention of the timex
    indiglo watch .
based on the late 1960's television show by the same name , the mod squad tells the tale of three
    reformed criminals under the employ of the police to go undercover .
however , things go wrong as evidence gets stolen and they are immediately under suspicion .
of course , the ads make it seem like so much more .
quick cuts , cool music , claire dane's nice hair and cute outfits , car
```

2.2   Input text from text file

NLTK corpora are very useful for testing NLP and text processing algorithms. But what if we wanted to load text from a text file? Let's try to load the text from the "alice.txt" file. First, copy the file "alice.txt" to your current working directory or retrieve the absolute path of the file. Then open the file, load its contents into a variable and close the file.

```python
f = open("alice.txt", "r") # Opens the file for reading only ("r")
text = f.read() # Store the contents of the file in variable "text". read() returns all the contents
    of the file
f.close() # Close the file
print(text) # Print the contents of variable "text"
```

The output will look like:

```
Alice was beginning to get very tired of sitting by her sister on the bank, and of having nothing to
    do: once or twice she had peeped into the book her sister was reading, but it had no pictures or
    conversations in it, "and what is the use of a book," thought Alice "without pictures or
    conversations?"

So she was considering in her own mind (as well as she could, for the hot day made her feel very
    sleepy and stupid), whether the pleasure of making a daisy-chain would be worth the trouble of
    getting up and picking the daisies, when suddenly a White Rabbit with pink eyes ran close by her.

There was nothing so very remarkable in that; nor did Alice think it so very much out of the way to
    hear the Rabbit say to itself, "Oh dear! Oh dear! I shall be late!" (when she thought it over
    afterwards, it occurred to her that she ought to have wondered at this, but at the time it all
    seemed quite natural); but when the Rabbit actually took a watch out of its waistcoat-pocket, and
    looked at it, and then hurried on, Alice started to her feet, for it flashed across her mind that
    she had never before seen a rabbit with either a waistcoat-pocket, or a watch to take out of it,
    and burning with curiosity, she ran across the field after it, and fortunately was just in time
    to see it pop down a large rabbit-hole under the hedge.
```

2.3   Text pre-processing

In the text output above, you can see that the text contains various sentences, a mix of lowercase and uppercase letters, there are some parentheses, as well as some punctuation marks.

2.3.1   Sentence Tokenisation

Let's tokenise the text, i.e. chop the text into pieces. In this case, the text is in the English language and punctuation marks are used to separate sentences from each other and a blank space is used to separate words from each other. For splitting a string into sentences, NLTK has the sent_tokenize() default tokeniser function.

Let's divide the text from alice.txt into sentences:

```python
from nltk import sent_tokenize # Import the sent_tokenize function from NLTK
sent_tokenize(text) # Tokenise "text" into sentences and print the output
```

The output will look like:

```
['Alice was beginning to get very tired of sitting by her sister on the bank, and of having nothing
    to do: once or twice she had peeped into the book her sister was reading, but it had no pictures
    or conversations in it, "and what is the use of a book," thought Alice "without pictures or
    conversations?"', 'So she was considering in her own mind (as well as she could, for the hot day
    made her feel very sleepy and stupid), whether the pleasure of making a daisy-chain would be
    worth the trouble of getting up and picking the daisies, when suddenly a White Rabbit with pink
    eyes ran close by her.', 'There was nothing so very remarkable in that; nor did Alice think it so
    very much out of the way to hear the Rabbit say to itself, "Oh dear!', 'Oh dear!', 'I shall be
    late!"', '(when she thought it over afterwards, it occurred to her that she ought to have
    wondered at this, but at the time it all seemed quite natural); but when the Rabbit actually took
    a watch out of its waistcoat-pocket, and looked at it, and then hurried on, Alice started to her
    feet, for it flashed across her mind that she had never before seen a rabbit with either a
    waistcoat-pocket, or a watch to take out of it, and burning with curiosity, she ran across the
    field after it, and fortunately was just in time to see it pop down a large rabbit-hole under the
    hedge.']
```

2.3.2   Word Tokenisation

Now let's divide the text from alice.txt into words using the word_tokenize() NLTK function:

```python
from nltk import word_tokenize # Import the word_tokenize function from NLTK
word_tokenize(text) # Tokenise "text" into words and print the output
```

The output will look like:

```
['Alice', 'was', 'beginning', 'to', 'get', 'very', 'tired', 'of', 'sitting', 'by', 'her', 'sister',
    'on', 'the', 'bank', ',', 'and', 'of', 'having', 'nothing', 'to', 'do', ':', 'once', 'or',
    'twice', 'she', 'had', 'peeped', 'into', 'the', 'book', 'her', 'sister', 'was', 'reading', ',',
    'but', 'it', 'had', 'no', 'pictures', 'or', 'conversations', 'in', 'it', ',', '``', 'and',
    'what', 'is', 'the', 'use', 'of', 'a', 'book', ',', "''", 'thought', 'Alice', '``', 'without',
    'pictures', 'or', 'conversations', '?', "''", 'So', 'she', 'was', 'considering', 'in', 'her',
    'own', 'mind', '(', 'as', 'well', 'as', 'she', 'could', ',', 'for', 'the', 'hot', 'day', 'made',
    'her', 'feel', 'very', 'sleepy', 'and', 'stupid', ')', ',', 'whether', 'the', 'pleasure', 'of',
    'making', 'a', 'daisy-chain', 'would', 'be', 'worth', 'the', 'trouble', 'of', 'getting', 'up',
    'and', 'picking', 'the', 'daisies', ',', 'when', 'suddenly', 'a', 'White', 'Rabbit', 'with',
    'pink', 'eyes', 'ran', 'close', 'by', 'her', '.', 'There', 'was', 'nothing', 'so', 'very',
    'remarkable', 'in', 'that', ';', 'nor', 'did', 'Alice', 'think', 'it', 'so', 'very', 'much',
    'out', 'of', 'the', 'way', 'to', 'hear', 'the', 'Rabbit', 'say', 'to', 'itself', ',', '``', 'Oh',
    'dear', '!', 'Oh', 'dear', '!', 'I', 'shall', 'be', 'late', '!', "''", '(', 'when', 'she',
    'thought', 'it', 'over', 'afterwards', ',', 'it', 'occurred', 'to', 'her', 'that', 'she',
    'ought', 'to', 'have', 'wondered', 'at', 'this', ',', 'but', 'at', 'the', 'time', 'it', 'all',
    'seemed', 'quite', 'natural', ')', ';', 'but', 'when', 'the', 'Rabbit', 'actually', 'took', 'a',
    'watch', 'out', 'of', 'its', 'waistcoat-pocket', ',', 'and', 'looked', 'at', 'it', ',', 'and',
    'then', 'hurried', 'on', ',', 'Alice', 'started', 'to', 'her', 'feet', ',', 'for', 'it',
    'flashed', 'across', 'her', 'mind', 'that', 'she', 'had', 'never', 'before', 'seen', 'a',
    'rabbit', 'with', 'either', 'a', 'waistcoat-pocket', ',', 'or', 'a', 'watch', 'to', 'take',
```

```
    'out', 'of', 'it', ',', 'and', 'burning', 'with', 'curiosity', ',', 'she', 'ran', 'across',
    'the', 'field', 'after', 'it', ',', 'and', 'fortunately', 'was', 'just', 'in', 'time', 'to',
    'see', 'it', 'pop', 'down', 'a', 'large', 'rabbit-hole', 'under', 'the', 'hedge', '.']
```

What if we wanted to tokenise each individual sentence of the text one by one?

```python
for sentence in sent_tokenize(text): # Iterate the sentences in "text"
    print(word_tokenize(sentence)) # Print the word tokenisation results. Don't forget the four
     spaces indentation that Python expects!
```

The output will look like:

```
['Alice', 'was', 'beginning', 'to', 'get', 'very', 'tired', 'of', 'sitting', 'by', 'her', 'sister',
    'on', 'the', 'bank', ',', 'and', 'of', 'having', 'nothing', 'to', 'do', ':', 'once', 'or',
    'twice', 'she', 'had', 'peeped', 'into', 'the', 'book', 'her', 'sister', 'was', 'reading', ',',
    'but', 'it', 'had', 'no', 'pictures', 'or', 'conversations', 'in', 'it', ',', '``', 'and',
    'what', 'is', 'the', 'use', 'of', 'a', 'book', ',', "'", 'thought', 'Alice', '``', 'without',
    'pictures', 'or', 'conversations', '?', "'"]
['So', 'she', 'was', 'considering', 'in', 'her', 'own', 'mind', '(', 'as', 'well', 'as', 'she',
    'could', ',', 'for', 'the', 'hot', 'day', 'made', 'her', 'feel', 'very', 'sleepy', 'and',
    'stupid', ')', ',', 'whether', 'the', 'pleasure', 'of', 'making', 'a', 'daisy-chain', 'would',
    'be', 'worth', 'the', 'trouble', 'of', 'getting', 'up', 'and', 'picking', 'the', 'daisies', ',',
    'when', 'suddenly', 'a', 'White', 'Rabbit', 'with', 'pink', 'eyes', 'ran', 'close', 'by', 'her',
    '.']
['There', 'was', 'nothing', 'so', 'very', 'remarkable', 'in', 'that', ';', 'nor', 'did', 'Alice',
    'think', 'it', 'so', 'very', 'much', 'out', 'of', 'the', 'way', 'to', 'hear', 'the', 'Rabbit',
    'say', 'to', 'itself', ',', '``', 'Oh', 'dear', '!']
['Oh', 'dear', '!']
['I', 'shall', 'be', 'late', '!', "'"]
['(', 'when', 'she', 'thought', 'it', 'over', 'afterwards', ',', 'it', 'occurred', 'to', 'her',
    'that', 'she', 'ought', 'to', 'have', 'wondered', 'at', 'this', ',', 'but', 'at', 'the', 'time',
    'it', 'all', 'seemed', 'quite', 'natural', ')', ';', 'but', 'when', 'the', 'Rabbit', 'actually',
    'took', 'a', 'watch', 'out', 'of', 'its', 'waistcoat-pocket', ',', 'and', 'looked', 'at', 'it',
    ',', 'and', 'then', 'hurried', 'on', ',', 'Alice', 'started', 'to', 'her', 'feet', ',', 'for',
    'it', 'flashed', 'across', 'her', 'mind', 'that', 'she', 'had', 'never', 'before', 'seen', 'a',
    'rabbit', 'with', 'either', 'a', 'waistcoat-pocket', ',', 'or', 'a', 'watch', 'to', 'take',
    'out', 'of', 'it', ',', 'and', 'burning', 'with', 'curiosity', ',', 'she', 'ran', 'across',
    'the', 'field', 'after', 'it', ',', 'and', 'fortunately', 'was', 'just', 'in', 'time', 'to',
    'see', 'it', 'pop', 'down', 'a', 'large', 'rabbit-hole', 'under', 'the', 'hedge', '.']
```

2.3.3 Lowercasing

As you can see in the list of tokens above, some words have uppercase letters. Let's convert all characters in our tokens to lowercase. We can use the lower() function to convert all characters in a string to lowercase. For example:

```python
test_string = "uNIverSIty"
test_string_lowercase = test_string.lower()
print(test_string_lowercase)
```

The output will look like:

```
university
```

Let's now convert all words in the text from alice.txt to lowercase, iterating through each sentence and word, and saving the lowercase words in a list:

```python
words_lowercase = []
for sentence in sent_tokenize(text): # Iterate the sentences in "text"
    for word in word_tokenize(sentence): # Iterate the words in "sentence"
        words_lowercase.append(word.lower()) # Convert word to lowercase and add it to the
     "words_lowercase" list
print(words_lowercase) # Print the list of lowercase words
```

The output will look like:

```
['alice', 'was', 'beginning', 'to', 'get', 'very', 'tired', 'of', 'sitting', 'by', 'her', 'sister',
    'on', 'the', 'bank', ',', 'and', 'of', 'having', 'nothing', 'to', 'do', ':', 'once', 'or',
    'twice', 'she', 'had', 'peeped', 'into', 'the', 'book', 'her', 'sister', 'was', 'reading', ',',
    'but', 'it', 'had', 'no', 'pictures', 'or', 'conversations', 'in', 'it', ',', '``', 'and',
    'what', 'is', 'the', 'use', 'of', 'a', 'book', ',', "'", 'thought', 'alice', '``', 'without',
    'pictures', 'or', 'conversations', '?', "'", 'so', 'she', 'was', 'considering', 'in', 'her',
    'own', 'mind', '(', 'as', 'well', 'as', 'she', 'could', ',', 'for', 'the', 'hot', 'day', 'made',
    'her', 'feel', 'very', 'sleepy', 'and', 'stupid', ')', ',', 'whether', 'the', 'pleasure', 'of',
    'making', 'a', 'daisy-chain', 'would', 'be', 'worth', 'the', 'trouble', 'of', 'getting', 'up',
    'and', 'picking', 'the', 'daisies', ',', 'when', 'suddenly', 'a', 'white', 'rabbit', 'with',
    'pink', 'eyes', 'ran', 'close', 'by', 'her', '.', 'there', 'was', 'nothing', 'so', 'very',
    'remarkable', 'in', 'that', ';', 'nor', 'did', 'alice', 'think', 'it', 'so', 'very', 'much',
    'out', 'of', 'the', 'way', 'to', 'hear', 'the', 'rabbit', 'say', 'to', 'itself', ',', '``', 'oh',
    'dear', '!', 'oh', 'dear', '!', 'i', 'shall', 'be', 'late', '!', "'", '(', 'when', 'she',
    'thought', 'it', 'over', 'afterwards', ',', 'it', 'occurred', 'to', 'her', 'that', 'she',
    'ought', 'to', 'have', 'wondered', 'at', 'this', ',', 'but', 'at', 'the', 'time', 'it', 'all',
    'seemed', 'quite', 'natural', ')', ';', 'but', 'when', 'the', 'rabbit', 'actually', 'took', 'a',
    'watch', 'out', 'of', 'its', 'waistcoat-pocket', ',', 'and', 'looked', 'at', 'it', ',', 'and',
    'then', 'hurried', 'on', ',', 'alice', 'started', 'to', 'her', 'feet', ',', 'for', 'it',
    'flashed', 'across', 'her', 'mind', 'that', 'she', 'had', 'never', 'before', 'seen', 'a',
    'rabbit', 'with', 'either', 'a', 'waistcoat-pocket', ',', 'or', 'a', 'watch', 'to', 'take',
    'out', 'of', 'it', ',', 'and', 'burning', 'with', 'curiosity', ',', 'she', 'ran', 'across',
    'the', 'field', 'after', 'it', ',', 'and', 'fortunately', 'was', 'just', 'in', 'time', 'to',
    'see', 'it', 'pop', 'down', 'a', 'large', 'rabbit-hole', 'under', 'the', 'hedge', '.']
```

As you can see, there are no uppercase characters left in the words list.

2.3.4   Stop words removal

Stop words are common words in any language that don't hold semantic meaning of their own. In many NLP applications, it is useful to remove stop words from a text. To achieve this, we use stop words lists that have been compiled for each language. We can access the list of English stop words from the NLTK library using the following:

```
nltk.download('stopwords') # Download the stopwords lists from NLTK

from nltk.corpus import stopwords # Import the stop words lists from NLTK
stopwords_english = stopwords.words('english') # Load the stop words list for English in variable
    "stopwords_english"
print(stopwords_english) # Print the "stopwords_english" list
```

The output will look like:

```
['i', 'me', 'my', 'myself', 'we', 'our', 'ours', 'ourselves', 'you', "you're", "you've", "you'll",
    "you'd", 'your', 'yours', 'yourself', 'yourselves', 'he', 'him', 'his', 'himself', 'she',
    "she's", 'her', 'hers', 'herself', 'it', "it's", 'its', 'itself', 'they', 'them', 'their',
    'theirs', 'themselves', 'what', 'which', 'who', 'whom', 'this', 'that', "that'll", 'these',
    'those', 'am', 'is', 'are', 'was', 'were', 'be', 'been', 'being', 'have', 'has', 'had', 'having',
    'do', 'does', 'did', 'doing', 'a', 'an', 'the', 'and', 'but', 'if', 'or', 'because', 'as',
    'until', 'while', 'of', 'at', 'by', 'for', 'with', 'about', 'against', 'between', 'into',
    'through', 'during', 'before', 'after', 'above', 'below', 'to', 'from', 'up', 'down', 'in',
    'out', 'on', 'off', 'over', 'under', 'again', 'further', 'then', 'once', 'here', 'there', 'when',
    'where', 'why', 'how', 'all', 'any', 'both', 'each', 'few', 'more', 'most', 'other', 'some',
    'such', 'no', 'nor', 'not', 'only', 'own', 'same', 'so', 'than', 'too', 'very', 's', 't', 'can',
    'will', 'just', 'don', "don't", 'should', "should've", 'now', 'd', 'll', 'm', 'o', 're', 've',
    'y', 'ain', 'aren', "aren't", 'couldn', "couldn't", 'didn', "didn't", 'doesn', "doesn't", 'hadn',
    "hadn't", 'hasn', "hasn't", 'haven', "haven't", 'isn', "isn't", 'ma', 'mightn', "mightn't",
    'mustn', "mustn't", 'needn', "needn't", 'shan', "shan't", 'shouldn', "shouldn't", 'wasn',
    "wasn't", 'weren', "weren't", 'won', "won't", 'wouldn', "wouldn't"]
```

As you can see, the list of words from alice.txt contains a number of stop words. Let's use NLTK's English stop words list to remove them:

```
words_lowercase_nostopwords = [] # Create empty list for remaining words
words_removed = [] # Create empty list for removed words
for word in words_lowercase: # Iterate lowercase words' list
    if word not in stopwords_english: # Check if word is in the stop words list
        words_lowercase_nostopwords.append(word)
    else:
        words_removed.append(word)
print(words_lowercase_nostopwords) # Print list of remaining words
print(words_removed) # Print list of removed words
```

The output will look like:

```
['alice', 'beginning', 'get', 'tired', 'sitting', 'sister', 'bank', ',', 'nothing', ':', 'twice',
    'peeped', 'book', 'sister', 'reading', ',', 'pictures', 'conversations', ',', '``', 'use',
    'book', ',', "'", 'thought', 'alice', '``', 'without', 'pictures', 'conversations', '?', "'",
    'considering', 'mind', '(', 'well', 'could', ',', 'hot', 'day', 'made', 'feel', 'sleepy',
    'stupid', ')', ',', 'whether', 'pleasure', 'making', 'daisy-chain', 'would', 'worth', 'trouble',
    'getting', 'picking', 'daisies', ',', 'suddenly', 'white', 'rabbit', 'pink', 'eyes', 'ran',
    'close', '.', 'nothing', 'remarkable', ';', 'alice', 'think', 'much', 'way', 'hear', 'rabbit',
    'say', ',', '``', 'oh', 'dear', '!', 'oh', 'dear', '!', 'shall', 'late', '!', "'", '(',
    'thought', 'afterwards', ',', 'occurred', 'ought', 'wondered', ',', 'time', 'seemed', 'quite',
    'natural', ')', ';', 'rabbit', 'actually', 'took', 'watch', 'waistcoat-pocket', ',', 'looked',
    ',', 'hurried', ',', 'alice', 'started', 'feet', ',', 'flashed', 'across', 'mind', 'never',
    'seen', 'rabbit', 'either', 'waistcoat-pocket', ',', 'watch', 'take', ',', 'burning',
    'curiosity', ',', 'ran', 'across', 'field', ',', 'fortunately', 'time', 'see', 'pop', 'large',
    'rabbit-hole', 'hedge', '.']
['was', 'to', 'very', 'of', 'by', 'her', 'on', 'the', 'and', 'of', 'having', 'to', 'do', 'once',
    'or', 'she', 'had', 'into', 'the', 'her', 'was', 'but', 'it', 'had', 'no', 'or', 'in', 'it',
    'and', 'what', 'is', 'the', 'of', 'a', 'or', 'so', 'she', 'was', 'in', 'her', 'own', 'as', 'as',
    'she', 'for', 'the', 'her', 'very', 'and', 'the', 'of', 'a', 'be', 'the', 'of', 'up', 'and',
    'the', 'when', 'a', 'with', 'by', 'her', 'there', 'was', 'so', 'very', 'in', 'that', 'nor',
    'did', 'it', 'so', 'very', 'out', 'of', 'the', 'to', 'the', 'to', 'itself', 'i', 'be', 'when',
    'she', 'it', 'over', 'it', 'to', 'her', 'that', 'she', 'to', 'have', 'at', 'this', 'but', 'at',
    'the', 'it', 'all', 'but', 'when', 'the', 'a', 'out', 'of', 'its', 'and', 'at', 'it', 'and',
    'then', 'on', 'to', 'her', 'for', 'it', 'her', 'that', 'she', 'had', 'before', 'a', 'with', 'a',
    'or', 'a', 'to', 'out', 'of', 'it', 'and', 'with', 'she', 'the', 'after', 'it', 'and', 'was',
    'just', 'in', 'to', 'it', 'down', 'a', 'under', 'the']
```

As you can see in the list of the remaining words and the list of the removed words, we removed multiple words from the alice.txt text, including "was", "to", "she", "it", "for", etc.

2.3.5 Punctuation removal

If you look at the list of the remaining words, you can see that some of its contents are not actual words, but they are punctuation marks instead. Let's remove these punctuation marks from the words list:

```
from string import punctuation # Import the punctuation marks string
print(punctuation)
print("Variable type: ", type(punctuation))
```

The output will look like:

```
!"#$%&'()*+,-./:;<=>?@[\]^_`{|}~
Variable type: <class 'str'>
```

As you can see, the variable "punctuation" contains all the potential punctuation marks. However, the variable is a string. To help us iterate and compare with the list of words, we will first convert the string "punctuation" to a list with one character (punctuation mark) per element.

```
punctuation_list = list(punctuation) # Convert punctuation to a list
print(punctuation_list)

words_lowercase_nostopwords_no_punctuation = []
for word in words_lowercase_nostopwords:
```

```
    if word not in punctuation_list:
        words_lowercase_nostopwords_no_punctuation.append(word)
print(words_lowercase_nostopwords_no_punctuation)
```

The output will look like:

```
['!', '"', '#', '$', '%', '&', "'", '(', ')', '*', '+', ',', '-', '.', '/', ':', ';', '<', '=', '>',
    '?', '@', '[', '\\', ']', '^', '_', '`', '{', '|', '}', '~']
['alice', 'beginning', 'get', 'tired', 'sitting', 'sister', 'bank', 'nothing', 'twice', 'peeped',
    'book', 'sister', 'reading', 'pictures', 'conversations', '``', 'use', 'book', "''", 'thought',
    'alice', '``', 'without', 'pictures', 'conversations', "''", 'considering', 'mind', 'well',
    'could', 'hot', 'day', 'made', 'feel', 'sleepy', 'stupid', 'whether', 'pleasure', 'making',
    'daisy-chain', 'would', 'worth', 'trouble', 'getting', 'picking', 'daisies', 'suddenly', 'white',
    'rabbit', 'pink', 'eyes', 'ran', 'close', 'nothing', 'remarkable', 'alice', 'think', 'much',
    'way', 'hear', 'rabbit', 'say', '``', 'oh', 'dear', 'oh', 'dear', 'shall', 'late', "''",
    'thought', 'afterwards', 'occurred', 'ought', 'wondered', 'time', 'seemed', 'quite', 'natural',
    'rabbit', 'actually', 'took', 'watch', 'waistcoat-pocket', 'looked', 'hurried', 'alice',
    'started', 'feet', 'flashed', 'across', 'mind', 'never', 'seen', 'rabbit', 'either',
    'waistcoat-pocket', 'watch', 'take', 'burning', 'curiosity', 'ran', 'across', 'field',
    'fortunately', 'time', 'see', 'pop', 'large', 'rabbit-hole', 'hedge']
```

As you can see, the majority of the punctuation marks where removed. But what about the remaining " ' " and " ' "? The punctuation marks list that we used did not contain double quotes and as a consequence, " ' " and " ' " were treated as valid words. We can address this issue by adding the respective double quote characters in the list of punctuation marks.

Note that there are alternative ways for removing punctuation marks or special characters from text, like for example the use of regular expressions.

2.3.6   Stemming

Consider the words "walk", "walks", "walking", and "walked". It is evident that all these words are different forms of the word "walk". We can use stemming to reduce each word to its respective stem , i.e. the core meaning-bearing unit of a word. NLTK comes with some common stemming algorithms. Let's use the Porter Stemming algorithm to reduce these four words to their stems:

```
from nltk.stem import PorterStemmer # Import the Porter stemmer from NLTK

porter = PorterStemmer() # Create a Porter stemmer object
for word in ['walk','walks','walking','walked']:
    print(word,"->",porter.stem(word))
```

The output will look like:

```
walk -> walk
walks -> walk
walking -> walk
walked -> walk
```

Let's now stem the words from the alice.txt text using the Porter stemmer:

```
words_stemmed = []
for word in words_lowercase_nostopwords_no_punctuation:
    words_stemmed.append(porter.stem(word))
print(words_stemmed)
```

The output will look like:

```
['alic', 'begin', 'get', 'tire', 'sit', 'sister', 'bank', 'noth', 'twice', 'peep', 'book', 'sister',
    'read', 'pictur', 'convers', 'use', 'book', 'thought', 'alic', 'without', 'pictur', 'convers',
    'consid', 'mind', 'well', 'could', 'hot', 'day', 'made', 'feel', 'sleepi', 'stupid', 'whether',
    'pleasur', 'make', 'daisy-chain', 'would', 'worth', 'troubl', 'get', 'pick', 'daisi', 'suddenli',
    'white', 'rabbit', 'pink', 'eye', 'ran', 'close', 'noth', 'remark', 'alic', 'think', 'much',
    'way', 'hear', 'rabbit', 'say', 'oh', 'dear', 'oh', 'dear', 'shall', 'late', 'thought',
```

```
'afterward', 'occur', 'ought', 'wonder', 'time', 'seem', 'quit', 'natur', 'rabbit', 'actual',
'took', 'watch', 'waistcoat-pocket', 'look', 'hurri', 'alic', 'start', 'feet', 'flash', 'across',
'mind', 'never', 'seen', 'rabbit', 'either', 'waistcoat-pocket', 'watch', 'take', 'burn',
'curios', 'ran', 'across', 'field', 'fortun', 'time', 'see', 'pop', 'larg', 'rabbit-hol', 'hedg']
```

As you can see, the Porter stemming algorithm reduced the words from alice.txt to their stems. However, it is evident that a lot of the stems do not correspond to real words from the English language.

2.3.7 Lemmatisation

Lemmatisation can address this issue by reducing each word to the respective dictionary headword. Let's lemmatise the words "walk", "walks", "walking", and "walked" using NLTK's WordNetLemmatizer:

```python
nltk.download('wordnet') # Download the WordNetLemmatizer package
from nltk.stem import WordNetLemmatizer # Import the WordNetLemmatizer

wnl = WordNetLemmatizer() # Create a WordNetLemmatizer object
for word in ['walk','walks','walking','walked']:
    print(word,"->",wnl.lemmatize(word))
```

The output will look like:

```
walk -> walk
walks -> walk
walking -> walking
walked -> walked
```

As you can see, the words "walk" and "walks" were converted to the lemma "walk", but "walking" and "walked" were not changed. The reason for this is that the WordNetLemmatizer considers by default all inputs as nouns, thus "walks" is considered as the plural form of "walk" and is converted to "walk", but "walking" and "walked" are valid lemmas and remain unchanged. To lemmatise all words to their base verb form, we must indicate that we are inputting verb forms:

```python
for word in ['walk','walks','walking','walked']:
    print(word,"->",wnl.lemmatize(word,pos='v'))
```

The output will look like:

```
walk -> walk
walks -> walk
walking -> walk
walked -> walk
```

Using the "pos" argument, the input words were handled as verb forms from the WordNetLemmatizer and returned the base form "walk" for all words. The options for "pos" are noun (n), verb (v), adverb (r), adjective (a). More information about the WordNetLemmatizer here: https://www.nltk.org/api/nltk.stem.wordnet.html

Let's lemmatise the text from alice.txt, treating the input as nouns and then as verbs:

```python
lemmas_noun = []
lemmas_verb = []
for word in words_lowercase_nostopwords_no_punctuation:
    lemmas_noun.append(wnl.lemmatize(word,pos='n'))
    lemmas_verb.append(wnl.lemmatize(word,pos='v'))
print(lemmas_noun)
print(lemmas_verb)
```

The output will look like:

```
['alice', 'beginning', 'get', 'tired', 'sitting', 'sister', 'bank', 'nothing', 'twice', 'peeped',
    'book', 'sister', 'reading', 'picture', 'conversation', 'use', 'book', 'thought', 'alice',
    'without', 'picture', 'conversation', 'considering', 'mind', 'well', 'could', 'hot', 'day',
    'made', 'feel', 'sleepy', 'stupid', 'whether', 'pleasure', 'making', 'daisy-chain', 'would',
    'worth', 'trouble', 'getting', 'picking', 'daisy', 'suddenly', 'white', 'rabbit', 'pink', 'eye',
```

```
        'ran', 'close', 'nothing', 'remarkable', 'alice', 'think', 'much', 'way', 'hear', 'rabbit',
        'say', 'oh', 'dear', 'oh', 'dear', 'shall', 'late', 'thought', 'afterwards', 'occurred', 'ought',
        'wondered', 'time', 'seemed', 'quite', 'natural', 'rabbit', 'actually', 'took', 'watch',
        'waistcoat-pocket', 'looked', 'hurried', 'alice', 'started', 'foot', 'flashed', 'across', 'mind',
        'never', 'seen', 'rabbit', 'either', 'waistcoat-pocket', 'watch', 'take', 'burning', 'curiosity',
        'ran', 'across', 'field', 'fortunately', 'time', 'see', 'pop', 'large', 'rabbit-hole', 'hedge']
['alice', 'begin', 'get', 'tire', 'sit', 'sister', 'bank', 'nothing', 'twice', 'peep', 'book',
        'sister', 'read', 'picture', 'conversations', 'use', 'book', 'think', 'alice', 'without',
        'picture', 'conversations', 'consider', 'mind', 'well', 'could', 'hot', 'day', 'make', 'feel',
        'sleepy', 'stupid', 'whether', 'pleasure', 'make', 'daisy-chain', 'would', 'worth', 'trouble',
        'get', 'pick', 'daisies', 'suddenly', 'white', 'rabbit', 'pink', 'eye', 'run', 'close',
        'nothing', 'remarkable', 'alice', 'think', 'much', 'way', 'hear', 'rabbit', 'say', 'oh', 'dear',
        'oh', 'dear', 'shall', 'late', 'think', 'afterwards', 'occur', 'ought', 'wonder', 'time', 'seem',
        'quite', 'natural', 'rabbit', 'actually', 'take', 'watch', 'waistcoat-pocket', 'look', 'hurry',
        'alice', 'start', 'feet', 'flash', 'across', 'mind', 'never', 'see', 'rabbit', 'either',
        'waistcoat-pocket', 'watch', 'take', 'burn', 'curiosity', 'run', 'across', 'field',
        'fortunately', 'time', 'see', 'pop', 'large', 'rabbit-hole', 'hedge']
```

However, this approach is not practical. Ideally, we would like to know what part of speech each word is and use the lemmatiser accordingly.

2.3.8   Part of Speech (POS) tagging

Part of Speech (POS) tagging is used to detect which part of speech each word in a sentence refers to. Let's use NLTK's POS tagging algorithm to assign POS tags to each of the words in the sentence "I had been a student here for a long time.":

```
nltk.download('averaged_perceptron_tagger')

from nltk import pos_tag

pos_tagged_sentence = pos_tag(word_tokenize('I had been a student here for a long time'))
print(pos_tagged_sentence)
```

The output will look like:

```
[('I', 'PRP'), ('had', 'VBD'), ('been', 'VBN'), ('a', 'DT'), ('student', 'NN'), ('here', 'RB'),
    ('for', 'IN'), ('a', 'DT'), ('long', 'JJ'), ('time', 'NN')]
```

As you can see, each word in the sentence has been annotated with a POS tags. These POS tags refer to the Penn Treebank POS tags (https://catalog.ldc.upenn.edu/docs/LDC95T7/cl93.html). However, the WordNetLemmatizer expects different tag names. To address this issue, we have to first convert the Penn Treebank POS tags to the format expected by the WordNetLemmatizer.

```
'''Convert Penn Treebank POS tags to WordNet'''
def penn_to_wordnet(penn_pos_tag):
    tag_dictionary = {'NN':'n', 'JJ':'a','VB':'v', 'RB':'r'}
    try:
        # If the first two characters of the Penn Treebank POS tag are in the ``tag_dictionary''
        return tag_dictionary[penn_pos_tag[:2]]
    except:
        return 'n' # Default to Noun if no mapping avalable.

lemmas = []
for word, tag in pos_tagged_sentence:
    lemmas.append(wnl.lemmatize(word.lower(), pos=penn_to_wordnet(tag)))

print('I have been a student here for a long time')
print(lemmas)
```

The output will look like:

```
I had been a student here for a long time
['i', 'have', 'be', 'a', 'student', 'here', 'for', 'a', 'long', 'time']
```

As you can see, the sentence was properly lemmatised using POS tagging to inform the lemmatiser about the part of speech that each word refers to.

Note: Please note that the conversion from Penn Treebank POS tags to WordNet format used here is a simplification. Full conversion tables should be used for better results. Also, the code above uses exception handling via the "try" and "except" statements. You can read more about exception handling in Python from here: `https://docs.python.org/3/tutorial/errors.html`

Let's now lemmatise the alice.txt text:

```
lemmas_alice = []

for sent in sent_tokenize(text): # Tokenise text into sentences
    pos_tagged_sentence_alice = pos_tag(word_tokenize(sent)) # Get POS tags for each sentence
    for word, tag in pos_tagged_sentence_alice: # Iterate though POS tagged words
        if word.lower() not in punctuation_list: # Ignore words that are punctuation marks
            lemmas_alice.append(wnl.lemmatize(word.lower(), pos=penn_to_wordnet(tag))) # Lemmatise word
print(lemmas_alice)
```

The output will look like:

```
['alice', 'be', 'begin', 'to', 'get', 'very', 'tired', 'of', 'sit', 'by', 'her', 'sister', 'on',
    'the', 'bank', 'and', 'of', 'have', 'nothing', 'to', 'do', 'once', 'or', 'twice', 'she', 'have',
    'peep', 'into', 'the', 'book', 'her', 'sister', 'be', 'read', 'but', 'it', 'have', 'no',
    'picture', 'or', 'conversation', 'in', 'it', '``', 'and', 'what', 'be', 'the', 'use', 'of', 'a',
    'book', "'", 'think', 'alice', '``', 'without', 'picture', 'or', 'conversation', "'", 'so',
    'she', 'be', 'consider', 'in', 'her', 'own', 'mind', 'as', 'well', 'a', 'she', 'could', 'for',
    'the', 'hot', 'day', 'make', 'her', 'feel', 'very', 'sleepy', 'and', 'stupid', 'whether', 'the',
    'pleasure', 'of', 'make', 'a', 'daisy-chain', 'would', 'be', 'worth', 'the', 'trouble', 'of',
    'get', 'up', 'and', 'pick', 'the', 'daisy', 'when', 'suddenly', 'a', 'white', 'rabbit', 'with',
    'pink', 'eye', 'run', 'close', 'by', 'her', 'there', 'be', 'nothing', 'so', 'very', 'remarkable',
    'in', 'that', 'nor', 'do', 'alice', 'think', 'it', 'so', 'very', 'much', 'out', 'of', 'the',
    'way', 'to', 'hear', 'the', 'rabbit', 'say', 'to', 'itself', '``', 'oh', 'dear', 'oh', 'dear',
    'i', 'shall', 'be', 'late', "'", 'when', 'she', 'think', 'it', 'over', 'afterwards', 'it',
    'occur', 'to', 'her', 'that', 'she', 'ought', 'to', 'have', 'wonder', 'at', 'this', 'but', 'at',
    'the', 'time', 'it', 'all', 'seem', 'quite', 'natural', 'but', 'when', 'the', 'rabbit',
    'actually', 'take', 'a', 'watch', 'out', 'of', 'it', 'waistcoat-pocket', 'and', 'look', 'at',
    'it', 'and', 'then', 'hurry', 'on', 'alice', 'start', 'to', 'her', 'foot', 'for', 'it', 'flash',
    'across', 'her', 'mind', 'that', 'she', 'have', 'never', 'before', 'see', 'a', 'rabbit', 'with',
    'either', 'a', 'waistcoat-pocket', 'or', 'a', 'watch', 'to', 'take', 'out', 'of', 'it', 'and',
    'burn', 'with', 'curiosity', 'she', 'run', 'across', 'the', 'field', 'after', 'it', 'and',
    'fortunately', 'be', 'just', 'in', 'time', 'to', 'see', 'it', 'pop', 'down', 'a', 'large',
    'rabbit-hole', 'under', 'the', 'hedge']
```

2.4 Exercises

Exercise 2.1 Section 2.3.5: Create a new punctuation marks list to address the issue of the remaining " ' " and " ' ".

Exercise 2.2 Section 2.3.5: Remove stop words and punctuation marks without iterating twice through all words.

Exercise 2.3 Load the text from dune.txt. Compute the number of words in the text, not including punctuation marks.

Exercise 2.4 Lemmatise the text from dune.txt and print a list of all the lemmas. Remember to convert all words to lowercase and to remove punctuation marks.

Exercise 2.5 Create a list of the unique lemmas in dune.txt, count their number and print the list and the number of lemmas.

Exercise 2.6 Create a custom function to divide English text into sentences without using NLTK or other tokenisers. Consider that a sentence ends when one of the following characters occurs: ".", "?", "!". Also remember to take into consideration the line change character "\n". Test your function on the text from dune.txt.

Workshop 3: Text representation

In this lab we will work on different ways to represent text.

3.1 Load corpus

We will use the Gutenberg corpus from NLTK. Let's first load the Gutenberg corpus:

```python
import nltk # Import the NLTK library
nltk.download('gutenberg') # Download the gutenberg corpus
from nltk.corpus import gutenberg # Import the gutenberg corpus from NLTK
gutenberg.fileids() # List file-ids in the corpus
```

The output will look like:

```
['austen-emma.txt',
 'austen-persuasion.txt',
 'austen-sense.txt',
 'bible-kjv.txt',
 'blake-poems.txt',
 'bryant-stories.txt',
 'burgess-busterbrown.txt',
 'carroll-alice.txt',
 'chesterton-ball.txt',
 'chesterton-brown.txt',
 'chesterton-thursday.txt',
 'edgeworth-parents.txt',
 'melville-moby_dick.txt',
 'milton-paradise.txt',
 'shakespeare-caesar.txt',
 'shakespeare-hamlet.txt',
 'shakespeare-macbeth.txt',
 'whitman-leaves.txt']
```

Let's compute some statistics for each document in the Gutenberg corpus, like the number of words, the number of sentences, and the number of characters:

```python
print("Chars\tWords\tSents\tFile")
for fileid in gutenberg.fileids(): # Iterate through files in corpus
    num_chars = len(gutenberg.raw(fileid))
    num_words = len(gutenberg.words(fileid))
    num_sents = len(gutenberg.sents(fileid))
    print("%7.0f\t%7.0f\t%7.0f\t%s" % (num_chars,num_words,num_sents,fileid))
```

The output will look like:

```
Chars      Words   Sents  File
 887071   192427    7752  austen-emma.txt
 466292    98171    3747  austen-persuasion.txt
 673022   141576    4999  austen-sense.txt
4332554  1010654   30103  bible-kjv.txt
  38153     8354     438  blake-poems.txt
```

```
 249439     55563     2863  bryant-stories.txt
  84663     18963     1054  burgess-busterbrown.txt
 144395     34110     1703  carroll-alice.txt
 457450     96996     4779  chesterton-ball.txt
 406629     86063     3806  chesterton-brown.txt
 320525     69213     3742  chesterton-thursday.txt
 935158    210663    10230  edgeworth-parents.txt
1242990    260819    10059  melville-moby_dick.txt
 468220     96825     1851  milton-paradise.txt
 112310     25833     2163  shakespeare-caesar.txt
 162881     37360     3106  shakespeare-hamlet.txt
 100351     23140     1907  shakespeare-macbeth.txt
 711215    154883     4250  whitman-leaves.txt
```

3.2 Vocabulary

As we can see above, each document consists of a few thousand words. But these words are not unique. Natural language consists of words that convey meaning and are re-used and combined to form different sentences. The set of unique words that are used in each document constitutes its vocabulary.

3.2.1 Words in corpus

Consider the text "The new table is red. The blue table is broken." Let's compute its vocabulary. First we should tokenise the text into words, convert them all to lowercase and remove punctuation marks:

```python
from nltk import word_tokenize # Import the word_tokenize function from NLTK
from string import punctuation

punctuation_list = list(punctuation) # Convert string with punctuation marks to list

text = "The new table is red. The blue table is broken."

text_tokens_processed = []
text_tokens = word_tokenize(text) # Tokenise the text into words
for token in text_tokens: # Iterate through the available tokens
    if token not in punctuation_list: # Omit tokens that are punctuation marks
        text_tokens_processed.append(token.lower()) # Add lowercase version of token to list

print("List of processed words:",text_tokens_processed)
```

The output will look like:

```
List of processed words: ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']
```

3.2.2 Unique words in corpus

As you can see in the list of words, the words "the", "table" and "is" appear two times each in the text. Let's now compute the vocabulary of this text, i.e. the list of unique words used in this text. To achieve this, we will use Python's **set** type. A set is similar to a list but allows only unique elements.

```python
vocabulary = set() # Create an empty set
for word in text_tokens_processed: # Iterate through available words
    vocabulary.add(word) # Add word to set

print("Vocabulary:",vocabulary)
print("Vocabulary size:",len(vocabulary))

vocabulary2 = set(text_tokens_processed)
print("\nVocabulary2:",vocabulary2)
print("Vocabulary2 size:",len(vocabulary2))
```

The output will look like:

```
Vocabulary: {'the', 'red', 'broken', 'blue', 'table', 'is', 'new'}
Vocabulary size: 7

Vocabulary2: {'the', 'red', 'broken', 'blue', 'table', 'is', 'new'}
Vocabulary2 size: 7
```

As you can see, the vocabulary used by the text "The new table is red. The blue table is broken." consists of the seven following words: is, table, blue, red, broken, new, the. Note that you don't have to add the contents of a list in a set one-by-one, as shown for variable "vocabulary2".

3.2.3 Vocabulary of multiple documents

Let's now compute the vocabulary for each document in the Gutenberg corpus:

```python
for fileid in gutenberg.fileids(): # Iterate through documents in corpus
    vocabulary_of_document = set() # Create empty set
    for word in gutenberg.words(fileid): # Iterate through words in document
        if word not in punctuation_list: # Omit tokens that are punctuation marks
            vocabulary_of_document.add(word.lower())
    print("%6.0f\t%s" % (len(vocabulary_of_document),fileid))
```

The output will look like:

```
  7328    austen-emma.txt
  5820    austen-persuasion.txt
  6388    austen-sense.txt
 12755    bible-kjv.txt
  1521    blake-poems.txt
  3925    bryant-stories.txt
  1547    burgess-busterbrown.txt
  2622    carroll-alice.txt
  8313    chesterton-ball.txt
  7780    chesterton-brown.txt
  6335    chesterton-thursday.txt
  8432    edgeworth-parents.txt
 17215    melville-moby_dick.txt
  9007    milton-paradise.txt
  3019    shakespeare-caesar.txt
  4703    shakespeare-hamlet.txt
  3451    shakespeare-macbeth.txt
 12437    whitman-leaves.txt
```

As you can see, we computed the vocabulary for each document in the Gutenberg corpus and printed its size. However, vocabularies from different documents are expected to have similar words in them since all texts are in the same language. Let's compute the vocabulary that covers all documents in the dataset:

```python
vocabulary_of_corpus = set() # Create empty set
for fileid in gutenberg.fileids(): # Iterate through documents in corpus
    for word in gutenberg.words(fileid): # Iterate through words in document
        if word not in punctuation_list: # Omit tokens that are punctuation marks
            vocabulary_of_corpus.add(word.lower())

print("Vocabulary of Gutenberg corpus:",len(vocabulary_of_corpus),"words")
```

The output will look like:

```
Vocabulary of Gutenberg corpus: 42314 words
```

As you can see, the vocabulary that covers all documents in the Gutenberg corpus is much larger than individual document vocabularies but significantly smaller that the sum of all the individual vocabularies.

3.3   One-hot encoding

3.3.1   One-hot encoding of words in vocabulary

Consider again the text "The new table is red. The blue table is broken." We have already computed its vocabulary and would like to compute the One-Hot representation of each word in the vocabulary.

```python
from numpy import array # Import array type from numpy
from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import OneHotEncoder

vocabulary = ['is', 'table', 'blue', 'red', 'broken', 'new', 'the']
data = array(vocabulary) # Convert to array because it is required by the LabelEncoder() object
print(data,"\n")

# Integer encoding - Assigns a unique index to each unique word
label_encoder = LabelEncoder()
integer_encoded = label_encoder.fit_transform(data)
print(integer_encoded,"\n")

# One-Hot encoding - Assigns a One-Hot binary representation to each word
onehot_encoder = OneHotEncoder(sparse_output=False)
integer_encoded = integer_encoded.reshape(len(integer_encoded), 1)
onehot_encoded = onehot_encoder.fit_transform(integer_encoded)
print(onehot_encoded,"\n")

for i in range(len(data)):
    print(onehot_encoded[i],"->",data[i])
```

The output will look like:

```
['is' 'table' 'blue' 'red' 'broken' 'new' 'the']

[2 5 0 4 1 3 6]

[[0. 0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1. 0.]
 [1. 0. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 1. 0. 0. 0.]
 [0. 0. 0. 0. 0. 0. 1.]]

[0. 0. 1. 0. 0. 0. 0.] -> is
[0. 0. 0. 0. 0. 1. 0.] -> table
[1. 0. 0. 0. 0. 0. 0.] -> blue
[0. 0. 0. 0. 1. 0. 0.] -> red
[0. 1. 0. 0. 0. 0. 0.] -> broken
[0. 0. 0. 1. 0. 0. 0.] -> new
[0. 0. 0. 0. 0. 0. 1.] -> the
```

Let's create a dictionary to easily encode our text and one-hot encode the word "red":

```python
dictionary = {}
for i in range(len(data)):
    dictionary[data[i]] = onehot_encoded[i]

print(dictionary)

print("\nred =",dictionary['red'])
```

The output will look like:

```
{'is': array([0., 0., 1., 0., 0., 0., 0.]), 'table': array([0., 0., 0., 0., 0., 1., 0.]), 'blue':
    array([1., 0., 0., 0., 0., 0., 0.]), 'red': array([0., 0., 0., 0., 1., 0., 0.]), 'broken':
    array([0., 1., 0., 0., 0., 0., 0.]), 'new': array([0., 0., 0., 1., 0., 0., 0.]), 'the':
```

```
    array([0., 0., 0., 0., 0., 0., 1.])}
```

```
red = [0. 0. 0. 0. 1. 0. 0.]
```

Consider the one-hot encoded word $(0, 0, 0, 0, 1, 0, 0,)$. How can we convert it back to its respective real word? Let's create a function to do this:

```python
def get_label_from_dictionary(dictionary,value):
    for word, one_hot in dictionary.items(): # Iterate all pairs (word, one-hot representation) in
     the dictionary
        if (one_hot == value).all(): # Compare equality between numpy arrays
            return word

print(get_label_from_dictionary(dictionary,[0,0,0,0,1,0,0]))
```

The output will look like:

```
red
```

3.3.2   One-hot encoding of text

Consider again the text "The new table is red. The blue table is broken." We have already computed its vocabulary and the one-hot representation of each word in the vocabulary. How can we one-hot encode the whole text? To do so, we have to perform a logical OR operation between the one-hot vectors of its constituent words.

```python
from numpy import logical_or # Import the element-wise logical OR function from numpy
from numpy import zeros # Import the zeros function from numpy

text = "The new table is red. The blue table is broken."
print("List of processed words:",text_tokens_processed)

result = zeros(len(dictionary)) # Start from zero-valued vector - Convert to numpy array
for word in text_tokens_processed: # Iterate words in text
    print(result.astype(int), "OR", dictionary[word],"= ",end='')
    result = logical_or(result,dictionary[word]) # Compute the element-wise logical or between the
     partial result and the one-hot representation of the word
    print(result.astype(int))

print("\nOne-Hot encoded text:",result.astype(int))
```

The output will look like:

```
List of processed words: ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']
[0 0 0 0 0 0 0] OR [0. 0. 0. 0. 0. 0. 1.] = [0 0 0 0 0 0 1]
[0 0 0 0 0 0 1] OR [0. 0. 0. 1. 0. 0. 0.] = [0 0 0 1 0 0 1]
[0 0 0 1 0 0 1] OR [0. 0. 0. 0. 0. 1. 0.] = [0 0 0 1 0 1 1]
[0 0 0 1 0 1 1] OR [0. 0. 1. 0. 0. 0. 0.] = [0 0 1 1 0 1 1]
[0 0 1 1 0 1 1] OR [0. 0. 0. 0. 1. 0. 0.] = [0 0 1 1 1 1 1]
[0 0 1 1 1 1 1] OR [0. 0. 0. 0. 0. 0. 1.] = [0 0 1 1 1 1 1]
[0 0 1 1 1 1 1] OR [1. 0. 0. 0. 0. 0. 0.] = [1 0 1 1 1 1 1]
[1 0 1 1 1 1 1] OR [0. 0. 0. 0. 0. 1. 0.] = [1 0 1 1 1 1 1]
[1 0 1 1 1 1 1] OR [0. 0. 1. 0. 0. 0. 0.] = [1 0 1 1 1 1 1]
[1 0 1 1 1 1 1] OR [0. 1. 0. 0. 0. 0. 0.] = [1 1 1 1 1 1 1]

One-Hot encoded text: [1 1 1 1 1 1 1]
```

Let's create a function for one-hot encoding text and do the same for the text "the broken table":

```python
def get_text_one_hot_encoding(words_list, dictionary):
    result = zeros(len(dictionary)) # Start from zero-valued vector - Convert to numpy array
    for word in words_list: # Iterate words in text
        result = logical_or(result,dictionary[word]) # Compute the element-wise logical or between the
     partial result and the one-hot representation of the word
```

```
    return result.astype(int)

words_list = ['the','broken','table']

print("\nOne-Hot encoded text:",get_text_one_hot_encoding(words_list,dictionary))
```

The output will look like:

```
One-Hot encoded text: [0 1 0 0 0 1 1]
```

3.4 Term Frequency (TF) representation

3.4.1 Compute TF of words in text

Let's now compute the term frequency of each word in the text "The new table is red. The blue table is broken."

```
from nltk import FreqDist # Import the FreqDist function from NLTK

text = "The new table is red. The blue table is broken."
text_tokens_processed = ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']
vocabulary = {'new', 'broken', 'the', 'blue', 'table', 'red', 'is'}

tf = FreqDist(text_tokens_processed) # Compute term frequency of words
print(tf,"\n")

vocabulary = sorted(vocabulary) # Sort alphabetiacally for better presentation
for word in vocabulary:
    print("%5.0f %s" % (tf[word],word))
```

The output will look like:

```
<FreqDist with 7 samples and 10 outcomes>

    1 blue
    1 broken
    2 is
    1 new
    1 red
    2 table
    2 the
```

3.4.2 TF representation of documents

Let's now compute the TF representation of the text "The new table is red. The blue table is broken." and the text "The new table is broken":

```
text_tf = []
for word in vocabulary:
    text_tf.append(tf[word])

print(text_tf,"->",text)

text2 = "The new table is broken"
text2_tokens_processed = ['the','new','table','is','broken']
tf2 = FreqDist(text2_tokens_processed) # Compute term frequency of words

text2_tf = []
for word in vocabulary:
    text2_tf.append(tf2[word])

print(text2_tf,"->",text2)
```

The output will look like:

```
[1, 1, 2, 1, 1, 2, 2] -> The new table is red. The blue table is broken.
[0, 1, 1, 1, 0, 1, 1] -> The new table is broken
```

3.5   Term Frequency - Inverse Document Frequency (TF-IDF)

Consider a corpus consisting of the three following documents:

1. "The new table is red. The blue table was broken."

2. "The new movie that we watched yesterday was terrible."

3. "We raised the red and blue flag yesterday."

3.5.1   Document Frequency (DF)

Let's compute the Document Frequency (DF) of each word in the above corpus. Remember that the DF of a word is equal to the number of documents in a corpus that the word appears in.

```python
# Create a list of the lists of lowercase words without punctuation marks for each document
texts_words_processed = []
texts_words_processed.append(['the','new','table','is','red','the','blue','table','was','broken'])
texts_words_processed.append(['the','new','movie','that','we','watched','yesterday','was','terrible'])
texts_words_processed.append(['we','raised','the','red','and','blue','flag','yesterday'])

print(texts_words_processed)

# Create the vocabulary
vocabulary_texts = set()
for doc in texts_words_processed:
    for word in doc:
        vocabulary_texts.add(word)

vocabulary_texts = sorted(vocabulary_texts) # Sort vocabulary alphabetically for better presentation

print("\nVocabulary:",vocabulary_texts)

DF = dict() # Create an empty dictionary
for word in vocabulary_texts: # Iterate through words in vocabulary
    cnt = 0
    for doc in texts_words_processed: # Iterate through documents
        if word in doc:
            cnt += 1 # cnt += 1 is equal to cnt = cnt + 1
    DF[word] = cnt

print("\nDocument frequencies:",DF)
```

The output will look like:

```
[['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'was', 'broken'], ['the', 'new',
    'movie', 'that', 'we', 'watched', 'yesterday', 'was', 'terrible'], ['we', 'raised', 'the', 'red',
    'and', 'blue', 'flag', 'yesterday']]

Vocabulary: ['and', 'blue', 'broken', 'flag', 'is', 'movie', 'new', 'raised', 'red', 'table',
    'terrible', 'that', 'the', 'was', 'watched', 'we', 'yesterday']

Document frequencies: {'and': 1, 'blue': 2, 'broken': 1, 'flag': 1, 'is': 1, 'movie': 1, 'new': 2,
    'raised': 1, 'red': 2, 'table': 1, 'terrible': 1, 'that': 1, 'the': 3, 'was': 2, 'watched': 1,
    'we': 2, 'yesterday': 2}
```

3.5.2   Inverse Document Frequency (IDF)

The Inverse Document Frequency (IDF) of a word is the logarithmically scaled inverse fraction of the documents that contain the word (obtained by dividing the total number of documents by the number of documents

containing the term, and then taking the logarithm of that quotient). IDF is defined as:

$$IDF(t, D) = \log\left(\frac{N}{DF(t, D)}\right) \tag{3.1}$$

where t is a word (term), D the corpus, and N the number of documents in the corpus.

Let's compute the IDF for the examined corpus:

```python
import math # Import math library

N = 3 # The corpus contains 3 documents

IDF = dict() # Create an empty dictionary

for word in vocabulary_texts: # Iterate through words in vocabulary
    IDF[word] = math.log( N / DF[word] ) # Compute IDF of word

print("IDF:",IDF)
```

The output will look like:

```
IDF: {'and': 1.0986122886681098, 'blue': 0.4054651081081644, 'broken': 1.0986122886681098, 'flag':
    1.0986122886681098, 'is': 1.0986122886681098, 'movie': 1.0986122886681098, 'new':
    0.4054651081081644, 'raised': 1.0986122886681098, 'red': 0.4054651081081644, 'table':
    1.0986122886681098, 'terrible': 1.0986122886681098, 'that': 1.0986122886681098, 'the': 0.0,
    'was': 0.4054651081081644, 'watched': 1.0986122886681098, 'we': 0.4054651081081644, 'yesterday':
    0.4054651081081644}
```

As you can see above, the more documents that a word appeared in, the lower the IDF for the word. Note that IDF is 0 for the word "the" that appeared in all documents (as a consequence of $N = DF \Rightarrow IDF = log1 = 0$)

3.5.3   Term Frequency - Inverse Document Frequency (TF-IDF)

Term Frequency - Inverse Document Frequency (TF-IDF) is defined as the product of TF and IDF:

$$TFIDF(t, d, D) = TF(t, d) \cdot IDF(t, D) \tag{3.2}$$

where d is a document of corpus D.

Let's compute the TF-IDF for each word in each document of the examined corpus

```python
from nltk import FreqDist # Import the FreqDist function from NLTK

TF = []
for doc in texts_words_processed: # Iterate though documents
    TF.append(FreqDist(doc)) # Compute word frequency

print(TF,"\n")

TFIDF = []
for tf_doc in TF:
    tfidf_doc = dict()
    for word in vocabulary_texts: # Iterate through words in vocabulary
        tfidf_doc[word] = tf_doc[word] * IDF[word] # Compute TF-IDF - tf_doc is of type FreqDist and
    will return 0 for words that don't exist
    TFIDF.append(tfidf_doc)

cnt = 0
for tfidf_doc in TFIDF:
    print("Text",cnt,"TF-IDF:",tfidf_doc,"\n")
    cnt += 1
```

The output will look like:

```
[FreqDist({'the': 2, 'table': 2, 'new': 1, 'is': 1, 'red': 1, 'blue': 1, 'was': 1, 'broken': 1}),
    FreqDist({'the': 1, 'new': 1, 'movie': 1, 'that': 1, 'we': 1, 'watched': 1, 'yesterday': 1,
    'was': 1, 'terrible': 1}), FreqDist({'we': 1, 'raised': 1, 'the': 1, 'red': 1, 'and': 1, 'blue':
    1, 'flag': 1, 'yesterday': 1})]

Text 0 TF-IDF: {'and': 0.0, 'blue': 0.4054651081081644, 'broken': 1.0986122886681098, 'flag': 0.0,
    'is': 1.0986122886681098, 'movie': 0.0, 'new': 0.4054651081081644, 'raised': 0.0, 'red':
    0.4054651081081644, 'table': 2.1972245773362196, 'terrible': 0.0, 'that': 0.0, 'the': 0.0, 'was':
    0.4054651081081644, 'watched': 0.0, 'we': 0.0, 'yesterday': 0.0}

Text 1 TF-IDF: {'and': 0.0, 'blue': 0.0, 'broken': 0.0, 'flag': 0.0, 'is': 0.0, 'movie':
    1.0986122886681098, 'new': 0.4054651081081644, 'raised': 0.0, 'red': 0.0, 'table': 0.0,
    'terrible': 1.0986122886681098, 'that': 1.0986122886681098, 'the': 0.0, 'was':
    0.4054651081081644, 'watched': 1.0986122886681098, 'we': 0.4054651081081644, 'yesterday':
    0.4054651081081644}

Text 2 TF-IDF: {'and': 1.0986122886681098, 'blue': 0.4054651081081644, 'broken': 0.0, 'flag':
    1.0986122886681098, 'is': 0.0, 'movie': 0.0, 'new': 0.0, 'raised': 1.0986122886681098, 'red':
    0.4054651081081644, 'table': 0.0, 'terrible': 0.0, 'that': 0.0, 'the': 0.0, 'was': 0.0,
    'watched': 0.0, 'we': 0.4054651081081644, 'yesterday': 0.4054651081081644}
```

3.6   Exercises

Exercise 3.1 Create a dictionary with the one-hot encoding of the vocabulary of the carroll-alice.txt file from the Gutenberg corpus.

Exercise 3.2 Based on the one-hot encoding of the vocabulary of the carroll-alice.txt from the Gutenberg corpus, one-hot encode the sentences "this is an old house", "this is a new house", and "he left his house" and compute their cosine distance.

Exercise 3.3 Using the vocabulary of the carroll-alice.txt document from the Gutenberg corpus, compute the TF-IDF representations and the respective cosine and euclidean distances of the documents in a corpus containing alice.txt and dune.txt.

Exercise 3.4 Use alice.txt to create a vocabulary and add to it the "unknown" word "<UNK>". Use this vocabulary to create the one-hot, the TF, the log normalised TF, and the TF-IDF representations of the alice.txt and the dune.txt documents.

Workshop 4: N-Grams

4.1 N-grams

We know that the probability of a sequence of words can be computed as:

$$P(w_1, w_2, w_3, ..., w_n) = P(w_1) \cdot P(w_2|w_1) \cdot P(w_3|w_1, w_2) \cdot ... \cdot P(w_n|w_1, w_2, w_3, ..., w_{n-1})$$

Unfortunately, we will never have enough data to compute the probability for any given word sequence. However, we can make some simplification assumptions and use N-grams to approximate the probabilities.

Let's compute various N-grams for the document "The new table is red. The blue table is broken.". Let's first load the document and compute the list of lowercase words, remove punctuation and compute the vocabulary.

```python
text = "The new table is red. The blue table is broken."
words_processed = ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']
vocabulary = set() # Create an empty set
for word in words_processed: # Iterate through available words
    vocabulary.add(word) # Add word to set

print("Document:",text)
print("Pre-processed words:",words_processed)
print("Document size:",len(words_processed))
print("Vocabulary:",vocabulary)
print("Vocabulary size:",len(vocabulary))
```

The output will look like:

```
Document: The new table is red. The blue table is broken.
Pre-processed words: ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']
Document size: 10
Vocabulary: {'red', 'new', 'the', 'blue', 'broken', 'is', 'table'}
Vocabulary size: 7
```

As you can see, the pre-processed document consists of 10 words and uses a vocabulary of 7 words.

4.2 Unigrams (1-Grams)

4.2.1 Compute unigrams

Unigrams (1-Grams) make the assumption that the probability of a word in a sequence of words depends only on the word itself (0th order Markovian assumption). Let's now compute the unigrams for the document "The new table is red. The blue table is broken.". Each unique word in the vocabulary constitutes a unigram of the document. Let's compute the counts of each unigram in our document.

```python
import nltk
from nltk import FreqDist # Import the FreqDist function from NLTK

tf = FreqDist(words_processed) # Compute term frequency of words
print(tf,"\n")

vocabulary = sorted(vocabulary) # Sort alphabetically for better presentation
unigrams = dict() # Create empty dictionary for unigrams
```

```
for word in vocabulary:
    unigrams[word] = tf[word]

print(unigrams)
```

The output will look like:

```
<FreqDist with 7 samples and 10 outcomes>

{'blue': 1, 'broken': 1, 'is': 2, 'new': 1, 'red': 1, 'table': 2, 'the': 2}
```

4.2.2   Unigram probability

The probability of a unigram for a word w_n is computed as:

$$P(w_n) = \frac{count(w_n)}{Total\,words} = \frac{count(w_n)}{\sum_{i=1}^{|V|} count(w_i)}$$

where w_n is a word, V the vocabulary, and $|V|$ the size of the vocabulary. Also, remember that when using unigrams, it is assumed that $P(w_n|w_{n-1}) \approx P(w_n)$

Let's now compute the probability of each word in the vocabulary:

```
total_words = len(words_processed) # Compute total words in corpus

unigram_probabilities = dict() # Create empty dictionary for unigram probabilities
for word in unigrams:
    unigram_probabilities[word] = unigrams[word] / total_words # Compute P(w_n)

print("Unigram probabilities:",unigram_probabilities)
```

The output will look like:

```
Unigram probabilities:
{'blue': 0.1, 'broken': 0.1, 'is': 0.2, 'new': 0.1, 'red': 0.1, 'table': 0.2, 'the': 0.2}
```

4.2.3   Sentence probability

Let's now compute the probability of the sentences "the new table is red" and "the black table" using the unigrams that we have computed.

$$P(\text{the new table is red}) \approx P(\text{the}) \cdot P(\text{new}) \cdot P(\text{table}) \cdot P(\text{is}) \cdot P(\text{red})$$

$$P(\text{the black table}) \approx P(\text{the}) \cdot P(\text{black}) \cdot P(\text{table})$$

Keep in mind that for words that don't exist in our corpus, the probability should be $P(\text{``unknown''}) = \frac{0}{Total\,words} = 0$

```
from collections import defaultdict

pw = defaultdict(lambda: 0, unigram_probabilities) # Create a dictionary that will return 0 for
    unknown words

print(pw,"\n")

p_text1 = pw["the"]*pw["new"]*pw["table"]*pw["is"]*pw["red"]
p_text2 = pw["the"]*pw["black"]*pw["table"]

print("P(the new table is red)= %f" % p_text1)
print("P(the black table)= %f" % p_text2)
```

The output will look like:

```
defaultdict(<function <lambda> at 0x0131CDB0>, {'blue': 0.1, 'broken': 0.1, 'is': 0.2, 'new': 0.1,
    'red': 0.1, 'table': 0.2, 'the': 0.2})

P(the new table is red)= 0.000080
P(the black table)= 0.000000
```

Note that we used the data structure "defaultdict" (https://docs.python.org/3/library/collections.html#collections.defaultdict) for storing the unigram probabilities. The reason for not using the default dictionary type of Python is that we need to set the probability to 0 for any unigram that is unknown and thus doesn't have a probability associated with it.

4.2.4 Smoothing

Notice that the probability for the sentence "the black table" is 0, as a result of the word "black" not existing in our corpus. However, this is a valid sentence in the English language. We will apply Add-λ smoothing in order to address the issue of zero-valued probabilities for unknown words.

$$P_{\text{Add-}\lambda}(w_n) = \frac{count(w_n) + \lambda}{\lambda|V| + \sum_{i=1}^{|V|} count(w_i)}$$

Let's now compute again the probability of the sentences "the new table is red" and "the black table" using the unigrams that we have computed and Add-λ smoothing for $\lambda = 0.001$. Remember that the probability of an unknown word when Add-λ smoothing is used will be:

$$P_{\text{Add-}\lambda}(\text{"unknown"}) = \frac{0 + \lambda}{\lambda|V| + \sum_{i=1}^{|V|} count(w_i)} = \frac{\lambda}{\lambda|V| + \sum_{i=1}^{|V|} count(w_i)}$$

```python
V = len(vocabulary) # Compute words in vocabulary
total_words = len(words_processed) # Compute total words in corpus
l = 0.001 # Define lambda for Add-lambda smoothing

p_unknown = (0 + l) / ( (l*V) + total_words) # Compute the probability of unknown words using
    add-lambda smoothing
print("P(unknown)=%f\n" % p_unknown)

unigram_probabilities_addl = dict() # Create empty dictionary for unigram probabilities
for word in unigrams:
    unigram_probabilities_addl[word] = (unigrams[word] + l) / (total_words + (l*V)) # Compute P(w_n)

print("Unigram probabilities (Add-lambda smoothing):\n",unigram_probabilities_addl,"\n")

plw = defaultdict(lambda: p_unknown, unigram_probabilities_addl) # Create a dictionary that will
    return p_unknown for unknown words

pl_text1 = plw["the"]*plw["new"]*plw["table"]*plw["is"]*plw["red"]
pl_text2 = plw["the"]*plw["black"]*plw["table"]

print("P(the new table is red)= %f" % pl_text1)
print("P(the black table)= %f" % pl_text2)
```

The output will look like:

```
P(unknown)=0.000100

Unigram probabilities (Add-lambda smoothing):
 {'blue': 0.10002997901468971, 'broken': 0.10002997901468971, 'is': 0.1999600279804137, 'new':
    0.10002997901468971, 'red': 0.10002997901468971, 'table': 0.1999600279804137, 'the':
    0.1999600279804137}

P(the new table is red)= 0.000080
P(the black table)= 0.000004
```

As you can see above, we can now compute the probability of word sequences that contain words that were not included in the corpus we used for creating our unigrams.

4.3   Bigrams (2-Grams)

4.3.1   Compute bigrams

Bigrams (2-Grams) make a 1st order Markovian assumption that the probability of a word in a sequence of words depends on the word and the previous word.

Let's now compute the bigrams for the examined corpus:

```python
from nltk.util import ngrams

text = "The new table is red. The blue table is broken."
words_processed = ['the', 'new', 'table', 'is', 'red', 'the', 'blue', 'table', 'is', 'broken']

bigrams = ngrams(words_processed,2) # Compute the bigrams in the text

bigrams_unique = set() # Create empty set for unique bigrams
for bigram in bigrams:
    print(bigram)
    bigrams_unique.add(bigram) # Add bigram to set

print("\nUnique bigrams:\n",bigrams_unique)
```

The output will look like:

```
('the', 'new')
('new', 'table')
('table', 'is')
('is', 'red')
('red', 'the')
('the', 'blue')
('blue', 'table')
('table', 'is')
('is', 'broken')

Unique bigrams:
 {('the', 'new'), ('is', 'broken'), ('red', 'the'), ('is', 'red'), ('blue', 'table'), ('the',
    'blue'), ('table', 'is'), ('new', 'table')}
```

4.3.2   Bigram probability

The probability of the bigram (w_{n-1}, w_n) is computed as:

$$P(w_n|w_{n-1}) = \frac{count(w_{n-1}, w_n)}{count(w_{n-1})}$$

Let's now compute the probability of each unique bigram in the examined corpus:

```python
bigrams = ngrams(words_processed,2) # Compute the bigrams in the text

bigram_freq = FreqDist(bigrams).items() # Compute frequency distribution for all the bigrams in the
    text

print(bigram_freq)
```

The output will look like:

```
dict_items([(('the', 'new'), 1), (('new', 'table'), 1), (('table', 'is'), 2), (('is', 'red'), 1),
    (('red', 'the'), 1), (('the', 'blue'), 1), (('blue', 'table'), 1), (('is', 'broken'), 1)])
```

4.3.3   Sentence probability

Let's now compute again the probability of the sentences "the new table is red" and "the black table" using the bigrams and the unigrams that we have computed. We will use the symbols $<s>$ and $</s>$ to indicate

the start and the end of a sentence respectively

$$P(< s > \text{the new table is red} < /s >) \approx$$

$$P(\text{the}| < s >) \cdot P(\text{new}|\text{the}) \cdot P(\text{table}|\text{new}) \cdot P(\text{is}|\text{table}) \cdot P(\text{red}|\text{is}) \cdot P(< /s > |\text{red}) \approx$$

$$\frac{count(< s >, \text{the})}{count(< s >)} \cdot \frac{count(\text{the}, \text{new})}{count(\text{the})} \cdot \frac{count(\text{new}, \text{table})}{count(\text{new})} \cdot \frac{count(\text{table}, \text{is})}{count(\text{table})} \cdot \frac{count(\text{is}, \text{red})}{count(\text{is})} \cdot \frac{count(\text{red}, < /s >)}{count(\text{red})}$$

$$P(< s > \text{the black table} < /s >) \approx P(\text{the}| < s >) \cdot P(\text{black}|\text{the}) \cdot P(\text{table}|\text{black}) \cdot P(< /s > |\text{table})$$

$$\approx \frac{count(< s >, \text{the})}{count(< s >)} \cdot \frac{count(\text{the}, \text{black})}{count(\text{the})} \cdot \frac{count(\text{black}, \text{table})}{count(\text{black})} \cdot \frac{count(\text{table}, < /s >)}{count(\text{table})}$$

Keep in mind that for bigrams that don't exist in our corpus, the probability should be $P(\text{"unknown"}) = 0$. Please note that the use of the tokens $< s >$ and $< /s >$ is optional!

```python
text = "The new table is red. The blue table is broken."
# Add tokens indicating the start and end of a sentence in the respective position
text2 = "<s> The new table is red. </s> <s> The blue table is broken. </s>"
words_processed = ['<s>','the', 'new', 'table', 'is', 'red','</s>','<s>','the', 'blue', 'table',
    'is', 'broken','</s>']

vocabulary = set() # Create an empty set
for word in words_processed: # Iterate through available words
    vocabulary.add(word) # Add word to set

tf = FreqDist(words_processed) # Compute term frequency of words

vocabulary = sorted(vocabulary) # Sort alphabetically for better presentation
ugf = dict() # Create empty dictionary for unigram counts
for word in vocabulary:
    ugf[word] = tf[word]

ugf = defaultdict(lambda: 0, ugf) # Create a dictionary that will return 0 for unknown unigrams
print("Unigram counts:",ugf,"\n")

bigrams = ngrams(words_processed,2) # Compute the bigrams in the text
bigram_freq = FreqDist(bigrams).items() # Compute frequency distribution for all the bigrams in the
    text
print("Bigram counts:",bigram_freq,"\n")

bgf = defaultdict(lambda: 0, bigram_freq) # Create a dictionary that will return 0 for unknown bigrams

def p_big(bigram, bigram_frequencies, unigram_frequencies): # Create function to compute bigram
    probability
    if(bigram_frequencies[bigram]==0):
        return 0
    else:
        return bigram_frequencies[bigram] / unigram_frequencies[bigram[0]]

p_text1 = p_big(('<s>','the'),bgf,ugf)*p_big(('the','new'),bgf,ugf)*p_big(('new','table'),bgf,ugf)
*p_big(('table','is'),bgf,ugf)*p_big(('is','red'),bgf,ugf)*p_big(('red','</s>'),bgf,ugf)

p_text2 = p_big(('<s>','the'),bgf,ugf)*p_big(('the','black'),bgf,ugf)*
p_big(('black','table'),bgf,ugf)*p_big(('table','</s>'),bgf,ugf)

print("P(<s> the new table is red </s>)= %f" % p_text1)
print("P(<s> the black table </s>)= %f" % p_text2)
```

The output will look like:

```
Unigram counts: {'</s>': 2, '<s>': 2, 'blue': 1, 'broken': 1, 'is': 2, 'new': 1, 'red': 1, 'table':
    2, 'the': 2}

Bigram counts: dict_items([(('<s>', 'the'), 2), (('the', 'new'), 1), (('new', 'table'), 1),
    (('table', 'is'), 2), (('is', 'red'), 1), (('red', '</s>'), 1), (('</s>', '<s>'), 1), (('the',
    'blue'), 1), (('blue', 'table'), 1), (('is', 'broken'), 1), (('broken', '</s>'), 1)])
```

```
P(<s> the new table is red </s>)= 0.250000
P(<s> the black table </s>)= 0.000000
```

4.3.4   Smoothing

Notice that the probability for the sentence "$< s >$ the black table $< /s >$" is 0, as a result of the bigrams (black,the), (table,black), and (black,$</s>$) not existing in our corpus. However, this is a valid sentence in the English language. We will apply Add-λ smoothing in order to address the issue of zero-valued probabilities for unknown bigrams.

$$P_{\text{Add-}\lambda}(w_n|w_{n-1}) = \frac{count(w_{n-1}, w_n) + \lambda}{\lambda|V| + count(w_{n-1})}$$

Let's now compute again the probability of the sentences "$< s >$ the new table is red $< /s >$" and "$< s >$ the black table $< /s >$" using the bigrams that we have computed and Add-λ smoothing for $\lambda = 0.01$.

```python
def pl_big(bigram, bigram_frequencies, unigram_frequencies,l): # Create function to compute bigram
    probability with add-lambda smoothing
    return (bigram_frequencies[bigram] + l) / ( (l*len(unigram_frequencies)) +
    unigram_frequencies[bigram[0]])

l = 0.01

pl_text1 = pl_big(('<s>','the'),bgf,ugf,l)*pl_big(('the','new'),bgf,ugf,l)*
pl_big(('new','table'),bgf,ugf,l)*pl_big(('table','is'),bgf,ugf,l)*pl_big(('is','red'),bgf,ugf,l)*
pl_big(('red','</s>'),bgf,ugf,l)

pl_text2 = pl_big(('<s>','the'),bgf,ugf,l)*pl_big(('the','black'),bgf,ugf,l)*
pl_big(('black','table'),bgf,ugf,l)*pl_big(('table','</s>'),bgf,ugf,l)

print("P(<s> the new table is red </s>)= %f" % pl_text1)
print("P(<s> the black table </s>)= %f" % pl_text2)
```

The output will look like:

```
P(<s> the new table is red </s>)= 0.185455
P(<s> the black table </s>)= 0.000002
```

4.4   Trigrams (3-Grams)

4.4.1   Compute trigrams

Trigrams (3-Grams) make a 2nd order Markovian assumption that the probability of a word in a sequence of words depends on the word and the previous two words.

Let's now compute the trigrams for the examined corpus, after adding the tokens "<s>" "<s>" and "</s>" "</s>" at the beginning and end of each sentence respectively.

```python
text = "<s> <s> The new table is red. </s> </s> <s> <s> The blue table is broken. </s> </s>"
words_processed = ['<s>','<s>','the', 'new', 'table', 'is', 'red','</s>','</s>','<s>','<s>', 'the',
    'blue', 'table', 'is', 'broken','</s>','</s>']

trigrams = ngrams(words_processed,3) # Compute the trigrams in the text

trigrams_unique = set() # Create empty set for unique trigrams
for trigram in trigrams:
    print(trigram)
    trigrams_unique.add(trigram) # Add trigram to set

print("\nUnique trigrams:\n",trigrams_unique)
```

The output will look like:

```
('<s>', '<s>', 'the')
('<s>', 'the', 'new')
('the', 'new', 'table')
('new', 'table', 'is')
('table', 'is', 'red')
('is', 'red', '</s>')
('red', '</s>', '</s>')
('</s>', '</s>', '<s>')
('</s>', '<s>', '<s>')
('<s>', '<s>', 'the')
('<s>', 'the', 'blue')
('the', 'blue', 'table')
('blue', 'table', 'is')
('table', 'is', 'broken')
('is', 'broken', '</s>')
('broken', '</s>', '</s>')

Unique trigrams:
 {('<s>', 'the', 'blue'), ('is', 'red', '</s>'), ('</s>', '</s>', '<s>'), ('new', 'table', 'is'),
    ('blue', 'table', 'is'), ('is', 'broken', '</s>'), ('broken', '</s>', '</s>'), ('table', 'is',
    'broken'), ('table', 'is', 'red'), ('<s>', '<s>', 'the'), ('red', '</s>', '</s>'), ('the',
    'blue', 'table'), ('</s>', '<s>', '<s>'), ('<s>', 'the', 'new'), ('the', 'new', 'table')}
```

4.4.2 Trigram probability

The probability of the trigram (w_{n-2}, w_{n-1}, w_n) is computed as:

$$P(w_n|w_{n-2}, w_{n-1}) = \frac{count(w_{n-2}, w_{n-1}, w_n)}{count(w_{n-2}, w_{n-1})}$$

Let's now compute the probability of each unique trigram in the examined corpus:

```
trigrams = ngrams(words_processed,3) # Compute the trigrams in the text

trigram_freq = FreqDist(trigrams).items() # Compute frequency distribution for all the trigrams in
    the text

print(trigram_freq)
```

The output will look like:

```
dict_items([(('<s>', '<s>', 'the'), 2), (('<s>', 'the', 'new'), 1), (('the', 'new', 'table'), 1),
    (('new', 'table', 'is'), 1), (('table', 'is', 'red'), 1), (('is', 'red', '</s>'), 1), (('red',
    '</s>', '</s>'), 1), (('</s>', '</s>', '<s>'), 1), (('</s>', '<s>', '<s>'), 1), (('<s>', 'the',
    'blue'), 1), (('the', 'blue', 'table'), 1), (('blue', 'table', 'is'), 1), (('table', 'is',
    'broken'), 1), (('is', 'broken', '</s>'), 1), (('broken', '</s>', '</s>'), 1)])
```

4.4.3 Sentence probability

Let's now compute again the probability of the sentences "$< s > < s >$ the new table is red $< /s > < /s >$" and "$< s > < s >$the black table $< /s > < /s >$" using the trigrams and the bigrams that we have computed.

$$P(< s >< s > \text{the new table is red} < /s >< /s >) \approx$$
$$P(\text{the}| < s >, < s >) \cdot P(\text{new}|, < s >, \text{the}) \cdot P(\text{table}|\text{the}, \text{new}) \cdot P(\text{is}|\text{new}, \text{table}) \cdot$$
$$P(\text{red}|\text{table}, \text{is}) \cdot P(< /s > |\text{is}, \text{red}) \cdot P(< /s > |\text{red}, < /s >) \approx$$
$$\frac{count(< s >, < s >, \text{the})}{count(< s >, < s >)} \cdot \frac{count(< s >, \text{the}, \text{new})}{count(< s >, \text{the})} \cdot \frac{count(\text{new}, \text{table}, \text{is})}{count(\text{new}, \text{table})} \cdot \frac{count(\text{table}, \text{is}, \text{red})}{count(\text{table}, \text{is})} \cdot$$
$$\frac{count(\text{is}, \text{red}, < /s >)}{count(\text{is}, \text{red})} \cdot \frac{count(\text{red}, < /s >, < /s >)}{count(\text{red}, < /s >)}$$

$$P(< \text{s} >< \text{s} > \text{the black table} < /\text{s} >< /\text{s} >) \approx$$

$$P(\text{the}| < \text{s} >, < \text{s} >) \cdot P(\text{black}| < \text{s} >, \text{the}) \cdot P(\text{table}|\text{the}, \text{black}) \cdot P(< /\text{s} > |\text{black}, \text{table}) \cdot$$

$$P(< /\text{s} > |\text{table}, < /\text{s} >) \approx$$

$$\frac{count(< \text{s} >, < \text{s} >, \text{the})}{count(< \text{s} >, < \text{s} >)} \cdot \frac{count(< \text{s} >, \text{the}, \text{black})}{count(< \text{s} >, \text{the})} \cdot \frac{count(\text{the}, \text{black}, \text{table})}{count(\text{the}, \text{black})} \cdot$$

$$\frac{count(\text{black}, \text{table}, < /\text{s} >)}{count(\text{black}, \text{table})} \cdot \frac{count(\text{table}, < /\text{s} >, < /\text{s} >)}{count(\text{table}, < /\text{s} >)}$$

Keep in mind that for trigrams that don't exist in our corpus, the probability should be $P(\text{"unknown"}) = 0$. Also, please note that the use of the $< \text{s} >$ and $< /\text{s} >$ tokens is optional!

```
text = "The new table is red. The blue table is broken."
# Add tokens indicating the start and end of a sentence in the respective position
text3 = "<s> <s> The new table is red. </s> </s> <s> <s> The blue table is broken. </s> </s>"
words_processed = ['<s>','<s>','the', 'new', 'table', 'is', 'red','</s>','</s>','<s>','<s>','the',
    'blue', 'table', 'is', 'broken','</s>','</s>']

bigrams = ngrams(words_processed,2) # Compute the bigrams in the text
bigram_freq = FreqDist(bigrams).items() # Compute frequency distribution for all the bigrams in the
    text
print("Bigram counts:",bigram_freq,"\n")


trigrams = ngrams(words_processed,3) # Compute the trigrams in the text
trigram_freq = FreqDist(trigrams).items() # Compute frequency distribution for all the trigrams in
    the text
print("Trigram counts:",trigram_freq,"\n")


bgf = defaultdict(lambda: 0, bigram_freq) # Create a dictionary that will return 0 for unknown bigrams
tgf = defaultdict(lambda: 0, trigram_freq) # Create a dictionary that will return 0 for unknown
    trigrams


def p_trig(trigram, trigram_frequencies, bigram_frequencies): # Create function to compute trigram
    probability
    if(trigram_frequencies[trigram]==0):
        return 0
    else:
        return trigram_frequencies[trigram] / bigram_frequencies[(trigram[0],trigram[1])]

p_text1 = p_trig(('<s>','<s>','the'),tgf,bgf)*p_trig(('<s>','the','new'),tgf,bgf)*
p_trig(('the','new','table'),tgf,bgf)*p_trig(('new','table','is'),tgf,bgf)*
p_trig(('table','is','red'),tgf,bgf)*p_trig(('is','red','</s>'),tgf,bgf)*
p_trig(('red','</s>','</s>'),tgf,bgf)

p_text2 = p_trig(('<s>','<s>','the'),tgf,bgf)*p_trig(('<s>','the','black'),tgf,bgf)*
p_trig(('the','black','table'),tgf,bgf)*p_trig(('black','table','</s>'),tgf,bgf)*
p_trig(('table','</s>','</s>'),tgf,bgf)

print("P(<s> <s> the new table is red </s> </s>)= %f" % p_text1)
print("P(<s> <s> the black table </s> </s>)= %f" % p_text2)
```

The output will look like:

```
Bigram counts: dict_items([(('<s>', '<s>'), 2), (('<s>', 'the'), 2), (('the', 'new'), 1), (('new',
    'table'), 1), (('table', 'is'), 2), (('is', 'red'), 1), (('red', '</s>'), 1), (('</s>', '</s>'),
    2), (('</s>', '<s>'), 1), (('the', 'blue'), 1), (('blue', 'table'), 1), (('is', 'broken'), 1),
    (('broken', '</s>'), 1)])

Trigram counts: dict_items([(('<s>', '<s>', 'the'), 2), (('<s>', 'the', 'new'), 1), (('the', 'new',
    'table'), 1), (('new', 'table', 'is'), 1), (('table', 'is', 'red'), 1), (('is', 'red', '</s>'),
    1), (('red', '</s>', '</s>'), 1), (('</s>', '</s>', '<s>'), 1), (('</s>', '<s>', '<s>'), 1),
    (('<s>', 'the', 'blue'), 1), (('the', 'blue', 'table'), 1), (('blue', 'table', 'is'), 1),
    (('table', 'is', 'broken'), 1), (('is', 'broken', '</s>'), 1), (('broken', '</s>', '</s>'), 1)])

P(<s> <s> the new table is red </s> </s>)= 0.250000
```

```
P(<s> <s> the black table </s> </s>)= 0.000000
```

4.4.4 Smoothing

Notice that the probability for the sentence "$< s > < s >$ the black table $< /s > < /s >$" is 0, as a result of the trigrams ($< s >$,the,black), (the,black,table), (black,table,$</s>$), and (table,$</s>$,$</s>$) not existing in our corpus. However, this is a valid sentence in the English language. We will apply Add-λ smoothing in order to address the issue of zero-valued probabilities for unknown trigrams.

$$P_{\text{Add-}\lambda}(w_n|w_{n-2}, w_{n-1}) = \frac{count(w_{n-2}, w_{n-1}, w_n) + \lambda}{\lambda|V| + count(w_{n-2}, w_{n-1})}$$

Let's now compute again the probability of the sentences "$< s > < s >$ the new table is red $< /s > < /s >$" and "$< s > < s >$ the black table $< /s > < /s >$" using the trigrams that we have computed and Add-λ smoothing for $\lambda = 0.001$.

```python
words_processed = ['<s>','<s>','the', 'new', 'table', 'is', 'red','</s>','</s>','<s>','<s>','the',
    'blue', 'table', 'is', 'broken','</s>','</s>']
vocabulary = set() # Create an empty set
for word in words_processed: # Iterate through available words
    vocabulary.add(word) # Add word to set

V = len(vocabulary) # Get size of vocabulary

def pl_trig(trigram, trigram_frequencies, bigram_frequencies,l,V): # Create function to compute
     trigram probability with add-lambda smoothing
    return (trigram_frequencies[trigram] + l) / ( (l*V) + bigram_frequencies[(trigram[0],trigram[1])])

l = 0.001

pl_text1 = pl_trig(('<s>','<s>','the'),tgf,bgf,l,V)*pl_trig(('<s>','the','new'),tgf,bgf,l,V)*
pl_trig(('the','new','table'),tgf,bgf,l,V)*pl_trig(('new','table','is'),tgf,bgf,l,V)*
pl_trig(('table','is','red'),tgf,bgf,l,V)*pl_trig(('is','red','</s>'),tgf,bgf,l,V)*
pl_trig(('red','</s>','</s>'),tgf,bgf,l,V)

pl_text2 = pl_trig(('<s>','<s>','the'),tgf,bgf,l,V)*pl_trig(('<s>','the','black'),tgf,bgf,l,V)*
pl_trig(('the','black','table'),tgf,bgf,l,V)*pl_trig(('black','table','</s>'),tgf,bgf,l,V)*
pl_trig(('table','</s>','</s>'),tgf,bgf,l,V)

print("P(<s> <s> the new table is red </s> </s>)= %f" % pl_text1)
print("P(<s> <s> the black table </s> </s>)= %f" % pl_text2)
```

The output will look like:

```
P(<s> <s> the new table is red </s> </s>)= 0.239523
P(<s> <s> the black table </s> </s>)= 0.000001
```

4.5 The number underflow issue

Let's use again the unigram model that we computed in Section 4.2.3 to compute the probability for the sentence "the new table is the broken blue table".

```python
text = "the new table is the broken blue table"
words_list = ['the','new','table','is','the','broken','blue','table']

print(pw,"\n")

p = 1
for word in words_list:
    p = p * pw[word]
    print("P(%s)=%f" % (word,pw[word]))

print("\nP(%s)=%f" % (text,p))
```

The output will look like:

```
defaultdict(<function <lambda> at 0x017334B0>, {'blue': 0.1, 'broken': 0.1, 'is': 0.2, 'new': 0.1,
    'red': 0.1, 'table': 0.2, 'the': 0.2, 'black': 0})

P(the)=0.200000
P(new)=0.100000
P(table)=0.200000
P(is)=0.200000
P(the)=0.200000
P(broken)=0.100000
P(blue)=0.100000
P(table)=0.200000

P(the new table is the broken blue table)=0.000000
```

As you can see above, the probability of sentence was computed as 0. But this is not correct. All the words in the sentence have a probability higher than 0. If you use a calculator to compute the probability $P(\text{the new table is the broken blue table}) = P(\text{the}) \cdot P(\text{new}) \cdot P(\text{table}) \cdot P(\text{is}) \cdot P(\text{the}) \cdot P(\text{broken}) \cdot P(\text{blue}) \cdot P(\text{table})$, the result will be 0.00000032. However, storing this number requires more precision that the float number type in Python supports and as a result it causes the number to underflow and return the value of 0. To avoid this problem, we typically compute probabilities in log space. Remember that in log space:

$$log\Big(P(A) \cdot P(B) \cdot P(C) \cdot ... \cdot P(Z)\Big) = log(P(A)) + log(P(B)) + log(P(C)) + ... + log(P(Z))$$

```python
import math # Import math library

text = "the new table is the broken blue table"
words_list = ['the','new','table','is','the','broken','blue','table']

logp = 0
for word in words_list:
    logp = logp + math.log(pw[word])
    print("log(P(%s))=%f" % (word,math.log(pw[word])))

print("\nlog(P(%s))=%f" % (text,logp))
```

The output will look like:

```
log(P(the))=-1.609438
log(P(new))=-2.302585
log(P(table))=-1.609438
log(P(is))=-1.609438
log(P(the))=-1.609438
log(P(broken))=-2.302585
log(P(blue))=-2.302585
log(P(table))=-1.609438

log(P(the new table is the broken blue table))=-14.954945
```

As long as all the probabilities are computed in log space, a higher log probability will denote a higher probability, since $a > b \implies log(a) > log(b)$. For example, from above: $P(\text{the}) > P(\text{new})$ $(0.2 > 0.1)$ and $log(P(\text{the})) > log(P(\text{new}))$ $(-1.609438 > -2.302585)$.

4.6   Exercises

Exercise 4.1 Create the function get_sentence_probability_unigram(words_list,unigram_frequencies), which given a sentence in the form of a list of words (words_list) and a dictionary with the frequencies of each unigram from a corpus (unigram_frequencies) will return the probability of the sentence based on the unigram language model. Use the function to compute the probability of the sentence "The passage in the castle" based on a unigram model trained on the dune.txt text. **Note:** Remember to address the number underflow issue.

Exercise 4.2 Use the carroll-alice.txt document from the NLTK Gutenberg corpus to train a unigram, a bigram, and a trigram model. For simplicity, do not use any tokens for the start and end of a sentence. Use these models to predict the next word in the following sentences:

    (a) we went for a _____

    (b) the food was _____

    (c) the weather was _____

    (d) yesterday she had _____

    (e) she was _____

    (f) since yesterday _____

Exercise 4.3 Create the function get_sentence_probability_bigram(words_list, unigram_frequencies, bigram_frequencies), which given a sentence in the form of a list of words (words_list), a dictionary with the frequencies of each unigram from a corpus (unigram_frequencies), and a dictionary with the frequencies of each bigram from a corpus (bigram_frequencies) will return the probability of the sentence based on the bigram language model. Use the function to compute the probability of the sentence "It was a warm night" based on a bigram model trained on the dune.txt text. **Note:** Remember to address the number underflow issue.

Workshop 5: Word embeddings

In this workshop we are going to create word embeddings using word-word co-occurrence matrices based on the context of each word.

5.1 Word context

To compute the word-word co-occurrence matrix we have to count the occurrences of each word from the vocabulary, within the context of each word. The context of a word can be set as a specific number of words prior and after the word, or the whole sentence, or the whole text, or even the whole corpus. Let's first compute the context of the word "i" in the text "*I like playing tennis. I enjoy sports. Do I enjoy tennis?*", in the form of a word list for a context size equal to one word before and one after the word "i".

5.1.1 Load text

Fist, let's tokenise the text to create the respective words list.

```python
from nltk import word_tokenize # Import the word_tokenize function from NLTK
from string import punctuation

punctuation_list = list(punctuation) # Convert punctuation to a list

text = "I like playing tennis. I enjoy sports. Do I enjoy tennis?"

tokens = word_tokenize(text.lower()) # Tokenise "text" into words

words_list = []
for word in tokens:
    if(word not in punctuation_list):
        words_list.append(word)

print(text,"->",words_list)
```

The output will look like:

```
I like playing tennis. I enjoy sports. Do I enjoy tennis? -> ['i', 'like', 'playing', 'tennis', 'i',
    'enjoy', 'sports', 'do', 'i', 'enjoy', 'tennis']
```

5.1.2 Compute context words

Then let's compute the words within the context of the word "i":

```python
context_size = 1
query_word = "i"
context = []
for i in range(len(words_list)): # Iterate through word list
    if(words_list[i] == query_word): # Check if word is the query word
        print("Found '%s' at position %.0f. Context:" % (query_word,i))
        for j in range(i-context_size,i+context_size+1): # Iterate through the context
            if( (j != i) and (j>=0) and (j<len(words_list)) ): # Ignore query word and non-valid word
    indexes
                context.append(words_list[j]) # Add word to context list
```

```
            print("[%.0f][%.0f] %s" % (i,j,words_list[j]))

print("\nContext of '%s' -> %s" % (query_word,context))
```

The output will look like:

```
Found 'i' at position 0. Context:
[0][1] like
Found 'i' at position 4. Context:
[4][3] tennis
[4][5] enjoy
Found 'i' at position 8. Context:
[8][7] do
[8][9] enjoy

Context of 'i' -> ['like', 'tennis', 'enjoy', 'do', 'enjoy']
```

Indeed, if you manually inspect the text, you will see that these four words appear within the context of the word "i" when the context is defined as the one previous word and the one after.

5.1.3   Other contexts

Let's now compute the words within the context of the words "i" and "enjoy" for a context size equal to n words prior and after each word, for $n = 1, 2, 3$. To avoid writing the same code multiple times, we will define a function for computing the context words.

```
def get_context(word, words_list,context_size):
    context = []
    for i in range(len(words_list)): # Iterate through word list
        if(words_list[i] == word): # Check if word is the query word
            for j in range(i-context_size,i+context_size+1): # Iterate through the context
                if( (j != i) and (j>=0) and (j<len(words_list)) ): # Ignore query word and non-valid
 word indexes
                    context.append(words_list[j]) # Add word to context list
    return context

print("\nContext (size=%.0f) of '%s' -> %s\n" % (1,"i",get_context("i", words_list,1)))
print("\nContext (size=%.0f) of '%s' -> %s\n" % (2,"i",get_context("i", words_list,2)))
print("\nContext (size=%.0f) of '%s' -> %s\n" % (3,"i",get_context("i", words_list,3)))
print("\nContext (size=%.0f) of '%s' -> %s\n" % (1,"enjoy",get_context("enjoy", words_list,1)))
print("\nContext (size=%.0f) of '%s' -> %s\n" % (2,"enjoy",get_context("enjoy", words_list,2)))
print("\nContext (size=%.0f) of '%s' -> %s\n" % (3,"enjoy",get_context("enjoy", words_list,3)))
```

The output will look like:

```
Context (size=1) of 'i' -> ['like', 'tennis', 'enjoy', 'do', 'enjoy']


Context (size=2) of 'i' -> ['like', 'playing', 'playing', 'tennis', 'enjoy', 'sports', 'sports',
    'do', 'enjoy', 'tennis']


Context (size=3) of 'i' -> ['like', 'playing', 'tennis', 'like', 'playing', 'tennis', 'enjoy',
    'sports', 'do', 'enjoy', 'sports', 'do', 'enjoy', 'tennis']


Context (size=1) of 'enjoy' -> ['i', 'sports', 'i', 'tennis']


Context (size=2) of 'enjoy' -> ['tennis', 'i', 'sports', 'do', 'do', 'i', 'tennis']


Context (size=3) of 'enjoy' -> ['playing', 'tennis', 'i', 'sports', 'do', 'i', 'sports', 'do', 'i',
    'tennis']
```

5.2 Word-word co-occurrence matrix

5.2.1 Word-word co-occurrence matrix (Context size = 1)

Let's now compute the word-word co-occurrence matrix for the text, for a context equal to one word before and one after the word.

```python
vocabulary = set(words_list) # Create vocabulary of unique words
vocabulary = sorted(list(vocabulary)) # Convert vocabulary to list to preserve ordering and sort it
    for better presentation
print("Vocabulary:",vocabulary,"\n")

context_size = 1

print("%7s" % "", end='')
for word in vocabulary:
    print("\t%7s" % word, end='')
print("\n")

for word in vocabulary:
    print("%7s" % word,end='')
    context = get_context(word, words_list,context_size)
    for context_word in vocabulary:
        print("\t%7.0f" % context.count(context_word),end='') # Prints the number of times that
     context_word appears in the context list
    print("\n")
```

The output will look like:

```
Vocabulary: ['do', 'enjoy', 'i', 'like', 'playing', 'sports', 'tennis']
```

	do	enjoy	i	like	playing	sports	tennis
do	0	0	1	0	0	1	0
enjoy	0	0	2	0	0	1	1
i	1	2	0	1	0	0	1
like	0	0	1	0	1	0	0
playing	0	0	0	1	0	0	1
sports	1	1	0	0	0	0	0
tennis	0	1	1	0	1	0	0

Note that we used the set type to create a vocabulary of unique words but we then converted it to a list. Sets in python do not support ordering of their contents, neither have they indexes assigned to each of their elements.

5.2.2 Word-word co-occurrence matrix (Context size = 2)

Let's now compute the word-word co-occurrence matrix for the text, for a context equal to two words before and two after the word.

```python
context_size = 2

print("%7s" % "", end='')
for word in vocabulary:
    print("\t%7s" % word, end='')
print("\n")

for word in vocabulary:
    print("%7s" % word,end='')
    context = get_context(word, words_list,context_size)
```

```
    for context_word in vocabulary:
        print("\t%7.0f" % context.count(context_word),end='')
    print("\n")
```

The output will look like:

	do	enjoy	i	like	playing	sports	tennis
do	0	2	1	0	0	1	0
enjoy	2	0	2	0	0	1	2
i	1	2	0	1	2	2	2
like	0	0	1	0	1	0	1
playing	0	0	2	1	0	0	1
sports	1	1	2	0	0	0	0
tennis	0	2	2	1	1	0	0

5.2.3   Compute word-word co-occurrence matrix as numpy array

Let's create a function that given a vocabulary, a text in the form of a word list, and the context size, will return a numpy array with the word-word co-occurrence matrix.

```
import numpy as np

def compute_word_word_matrix(vocabulary,words_list,context_size):
    word_word_matrix = np.zeros(( len(vocabulary),len(vocabulary) ), dtype=int) # Create empty array
     of size VxV
    for i in range(len(vocabulary)):
        context = get_context(vocabulary[i], words_list,context_size)
        for j in range(len(vocabulary)):
            word_word_matrix[i,j] = context.count(vocabulary[j])
    return word_word_matrix


context_size = 2

word_word_matrix = compute_word_word_matrix(vocabulary,words_list,context_size)

print(word_word_matrix)
```

The output will look like:

```
[[0 2 1 0 0 1 0]
 [2 0 2 0 0 1 2]
 [1 2 0 1 2 2 2]
 [0 0 1 0 1 0 1]
 [0 0 2 1 0 0 1]
 [1 1 2 0 0 0 0]
 [0 2 2 1 1 0 0]]
```

5.2.4   Word-word co-occurrence matrix visualisation

Let's visualise the word-word co-occurrence matrix as a heatmap.

```
import matplotlib
import matplotlib.pyplot as plt

fig, ax = plt.subplots()
im = ax.imshow(word_word_matrix, cmap='viridis') # Create heatmap using the 'viridis' colour map
```

```python
# Show all ticks
ax.set_xticks(np.arange(len(vocabulary)))
ax.set_yticks(np.arange(len(vocabulary)))
# Label ticks with the respective list entries
ax.set_xticklabels(vocabulary)
ax.set_yticklabels(vocabulary)

# Rotate the tick labels and set their alignment.
plt.setp(ax.get_xticklabels(), rotation=45, ha="right",rotation_mode="anchor")

# Loop over data dimensions and create text annotations.
for i in range(len(vocabulary)):
    for j in range(len(vocabulary)):
        text = ax.text(j, i, word_word_matrix[i, j], ha="center", va="center", color="w")

ax.set_title("Word-word co-occurrence matrix (Context size=%.0f)" % context_size)

plt.colorbar(im) # Add colour bar with colour range
plt.show() # Show plot
```

The output will look like:
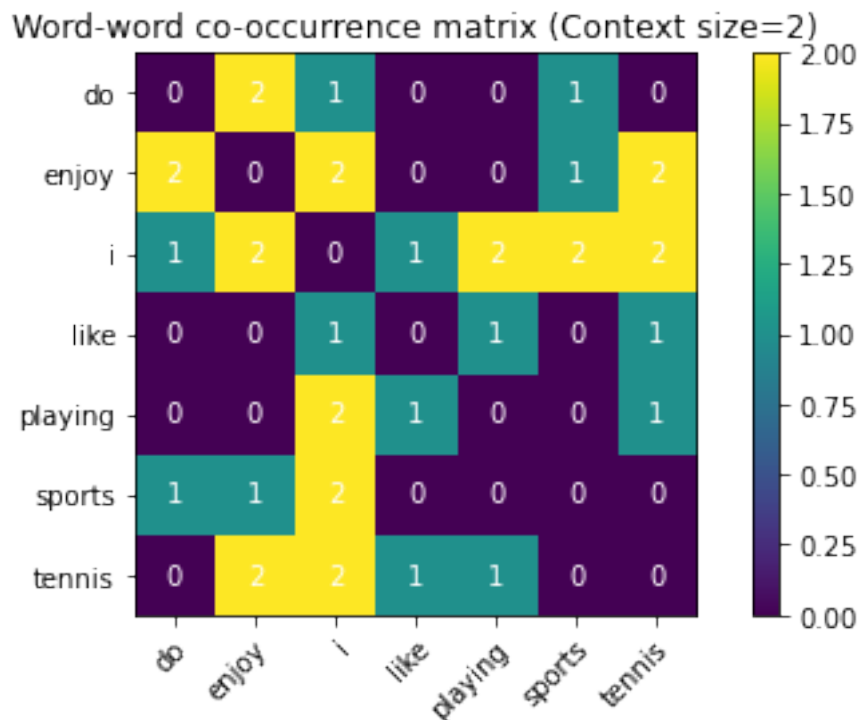
5.3 Word embeddings

5.3.1 Word embeddings computation

We will use the word-word co-occurrence matrix that we have computed to create the word embeddings for the words in our text's vocabulary.

```python
def get_word_embedding(word,word_word_matrix,vocabulary):
    word_index = vocabulary.index(word) # Gets word's index. Vocabulary must be of list type
    return word_word_matrix[word_index,:] # Return the word_index-th row of the word-word matrix

word_vectors = dict()
for word in vocabulary:
    word_vectors[word] = get_word_embedding(word,word_word_matrix,vocabulary)
```

```
    print(word,"->",get_word_embedding(word,word_word_matrix,vocabulary))

print("\n%s" % word_vectors)
```

The output will look like:

```
do -> [0 2 1 0 0 1 0]
enjoy -> [2 0 2 0 0 1 2]
i -> [1 2 0 1 2 2 2]
like -> [0 0 1 0 1 0 1]
playing -> [0 0 2 1 0 0 1]
sports -> [1 1 2 0 0 0 0]
tennis -> [0 2 2 1 1 0 0]

{'do': array([0, 2, 1, 0, 0, 1, 0]), 'enjoy': array([2, 0, 2, 0, 0, 1, 2]), 'i': array([1, 2, 0, 1,
    2, 2, 2]), 'like': array([0, 0, 1, 0, 1, 0, 1]), 'playing': array([0, 0, 2, 1, 0, 0, 1]),
    'sports': array([1, 1, 2, 0, 0, 0, 0]), 'tennis': array([0, 2, 2, 1, 1, 0, 0])}
```

5.3.2   Word embeddings visualisation

We can visualise word embeddings as vectors in a V-dimensional space, where V is the size of the vocabulary.
Let's visualise the vectors for the words "i" and "tennis" in the "enjoy" and "do" dimensions.

```
index_enjoy = vocabulary.index("enjoy") # Get index of word "enjoy" in vocabulary
index_do = vocabulary.index("do") # Get index of word "do" in vocabulary

# Create word embedding using only the values for the dimensions "enjoy" and "do"
embedding_i = word_vectors["i"][[index_enjoy,index_do]]
print("i ->",embedding_i)
embedding_tennis = word_vectors["tennis"][[index_enjoy,index_do]]
print("tennis ->",embedding_tennis,"\n")

fig = plt.subplots()
plt.plot([0,embedding_tennis[0]], [0,embedding_tennis[1]], 'g', label="tennis") # Plot line from
    (0,0) to the "tennis" coordinates
plt.plot([0,embedding_i[0]], [0,embedding_i[1]], 'b', label="i") # Plot line from (0,0) to the "i"
    coordinates
plt.xlabel('enjoy') # Set label for x axis
plt.ylabel('do') # Set label for y axis
plt.legend(loc="upper left") # Show plot legend at upper left location
plt.show() # Show plot
```

The output will look like:

```
i -> [2 1]
tennis -> [2 0]
```

5.3.3 Word embeddings distance

Let's now compute the pairwise cosine distance for all the words in the vocabulary using the word embeddings that we computed before.

```python
from scipy.spatial import distance

print("Words cosine distance:")
for word1 in vocabulary:
    for word2 in vocabulary:
        print(word1,"->",word2,"=",distance.cosine(word_vectors[word1],word_vectors[word2]))
```

The output will look like:

```
Words cosine distance:
do -> do = 0.0
do -> enjoy = 0.6603168897566213
do -> i = 0.42264973081037427
do -> like = 0.7642977396044841
do -> playing = 0.6666666666666667
do -> sports = 0.33333333333333337
do -> tennis = 0.2254033307585167
enjoy -> do = 0.6603168897566213
enjoy -> enjoy = 0.0
enjoy -> i = 0.4770236396315093
enjoy -> like = 0.3594873847796515
enjoy -> playing = 0.32063377951324257
enjoy -> sports = 0.32063377951324257
enjoy -> tennis = 0.6491767922771884
i -> do = 0.42264973081037427
i -> enjoy = 0.4770236396315093
i -> i = 0.0
i -> like = 0.45566894604818275
i -> playing = 0.7113248654051871
i -> sports = 0.7113248654051871
i -> tennis = 0.47825080525004915
like -> do = 0.7642977396044841
like -> enjoy = 0.3594873847796515
like -> i = 0.45566894604818275
like -> like = 0.0
like -> playing = 0.2928932188134524
like -> sports = 0.5285954792089682
like -> tennis = 0.4522774424948339
playing -> do = 0.6666666666666667
playing -> enjoy = 0.32063377951324257
playing -> i = 0.7113248654051871
playing -> like = 0.2928932188134524
playing -> playing = 0.0
playing -> sports = 0.33333333333333337
playing -> tennis = 0.3545027756320972
sports -> do = 0.33333333333333337
sports -> enjoy = 0.32063377951324257
sports -> i = 0.7113248654051871
sports -> like = 0.5285954792089682
sports -> playing = 0.33333333333333337
sports -> sports = 0.0
sports -> tennis = 0.2254033307585167
tennis -> do = 0.2254033307585167
tennis -> enjoy = 0.6491767922771884
tennis -> i = 0.47825080525004915
tennis -> like = 0.4522774424948339
tennis -> playing = 0.3545027756320972
tennis -> sports = 0.2254033307585167
tennis -> tennis = 0.0
```

5.4 Exercises

Exercise 5.1 Create the word-word co-occurrence matrix for the dune.txt text for a context size equal to the 3 words prior and after a word. Visualise the word-word co-occurrence matrix as a heatmap.

Exercise 5.2 Use the word-word co-occurrence matrix that you computed in Exercise 5.1 in order to compute the respective word embeddings for all words in the vocabulary of dune.txt. Then compute the pairwise cosine distance between all words in the vocabulary and visualise them as a heatmap.

Exercise 5.3 Create the word-word co-occurrence matrix for the gatsby.txt document, compute the respective word embeddings, compute the pairwise cosine distance for all words in the vocabulary, and visualise these distances as a heatmap. Then, select the 5 dimensions (words) with the most counts and compute the word embeddings for the vocabulary using only these 5 dimensions. Consider the context as the 10 words prior and after a word.

Workshop 6: Document embeddings for machine learning

In this workshop, we are going to use word embeddings in order to create document embeddings that will be used for training machine learning models for the task of classification.

6.1 Loading data with pandas

Pandas is a fast, powerful, flexible and easy to use open source data analysis and manipulation package, built on top of the Python programming language. You can find more information about the pandas Python package from here: `https://pandas.pydata.org/`

6.1.1 Load dataset file

Let's load the file testvectors.csv, which contains the word embeddings (of size 300) for a subset of the words from a 2017 dump of the English Wikipedia. The embeddings were created using the skipgram model and are stored in a comma-separated file, where each row contains the embedding of a word in 301 columns. The first column contains the word and the rest 300 columns the respective embedding.

We will use pandas to load the testvectors.csv file and store its content to a pandas dataframe object.

```python
import pandas as pd # Import the pandas library

df = pd.read_csv('testvectors.csv', header=None) # Read csv file. Indicate that there is no row with
    column titles
print( df.head(5), "\n") # Print the first 5 rows of the dataframe

count_row = df.shape[0] # Gives number of rows
count_col = df.shape[1] # Gives number of columns

print("\nTotal words:",count_row)
print("Dimensions:",count_col-1) # Subtract 1 for the word column
```

The output will look like:

```
       0         1         2         3         4         5         6  \
0    one  0.073525 -0.031703 0.054010 -0.040015 -0.011894 0.002958
1   time  0.061892 0.066106 0.026482 -0.122901 0.016603 0.024152
2  would -0.005455 -0.064055 0.106359 0.000271 -0.005658 0.017313
3   made -0.033224 -0.046773 -0.022644 -0.115277 -0.037984 0.109500
4   well  0.032183 0.058166 0.102063 -0.054714 -0.037364 0.013909


         7         8         9  ...       291       292       293       294  \
0 -0.065406 0.079166 0.099932 ...  0.048875 0.014754 -0.038729 0.033155
1 -0.057542 0.119501 0.033247 ... -0.043285 0.045578 -0.139174 0.109938
2 -0.030339 0.014298 0.030399 ... -0.029693 0.083669 -0.089212 0.063672
3 -0.138488 0.038567 0.001758 ...  0.005144 0.002312 -0.043491 0.091255
4  0.025827 0.072022 -0.050159 ...  0.017727 0.056166 -0.106694 0.046184


       295       296       297       298       299       300
0 -0.120390 -0.065746 -0.023745 0.012824 0.005162 -0.130008
```

```
1 -0.057847 0.010336 0.114048 0.042011 0.032165 0.051062
2 -0.048096 0.028711 0.032499 0.104135 0.001100 0.024307
3 -0.065307 -0.060637 -0.029168 0.034441 0.024149 0.003539
4 -0.091665 -0.088526 0.016668 0.061213 -0.015307 -0.076759


[5 rows x 301 columns]


Total words: 155
Dimensions: 300
```

6.1.2   Available words in dataset

Let's create a list with the available words from the word embeddings dataset from testvectors.csv.

```
available_words = df[0].tolist() # Convert the first column of the dataframe to a list and store to a
    variable

print(available_words)
```

The output will look like:

```
['one', 'time', 'would', 'made', 'well', 'family', 'use', 'took', 'could', 'home', 'served', 'large',
    'like', 'day', 'final', 'near', 'much', 'book', 'came', 'late', 'side', 'started', 'way', 'take',
    'without', 'old', 'making', 'field', 'never', 'across', 'see', 'features', 'seen', 'mother',
    'either', 'get', 'close', 'reached', 'white', 'change', 'female', 'beginning', 'allowed',
    'night', 'week', 'natural', 'ran', 'thought', 'woman', 'room', 'nearly', 'sister', 'acquired',
    'whether', 'ancient', 'actually', 'feet', 'bank', 'floor', 'occurred', 'stone', 'twice', 'visit',
    'say', 'quite', 'castle', 'think', 'pop', 'shape', 'getting', 'reading', 'nothing', 'boy',
    'standing', 'mind', 'ahead', 'weather', 'let', 'door', 'feel', 'step', 'eyes', 'hot', 'hair',
    'moment', 'worth', 'afterwards', 'departure', 'shall', 'lay', 'passage', 'watch', 'looked',
    'seemed', 'bed', 'sitting', 'pictures', 'feeling', 'hear', 'generations', 'trouble', 'warm',
    'suddenly', 'considering', 'burning', 'remarkable', 'bore', 'pink', 'hanging', 'pleasure',
    'shadow', 'peer', 'darkness', 'picking', 'tired', 'lamp', 'witch', 'conversations', 'ought',
    'pile', 'rabbit', 'hedge', 'curiosity', 'stupid', 'jewels', 'dear', 'twenty-six', 'vaulted',
    'awakened', 'wondered', 'bulky', 'frenzy', 'hooded', 'hurried', 'oh', 'fortunately',
    'unbearable', 'sleepy', 'flashed', 'glittering', 'in', 'by', 'matted', 'the', 'dimmed', 'of',
    'and', 'daisies', 'a', 'to', 'scurrying', 'for', 'crone', 'on', 'is']
```

6.1.3   Access specific row in pandas dataframe

Let's now attempt to access a specific row from the pandas dataframe that we have stored our embeddings in. To access the first row from the dataframe we have to do the following:

```
vector = df.iloc[0] # Access the 0-th row of the dataframe
print(vector)
```

The output will look like:

```
0          one
1      0.073525
2     -0.031703
3       0.05401
4     -0.040015
         ...
296   -0.065746
297   -0.023745
298    0.012824
299    0.005162
300   -0.130008
Name: 0, Length: 301, dtype: object
```

6.1.4 Access specific columns in pandas dataframe

Let's now attempt to access only specific columns from a row in the pandas dataframe. The first column of each row (column 0) contains the word, while the respective embedding is stored in columns 1-301. To access from column 1 to column 5 from the first row in the dataframe and from column 1 to 301, we have to do the following:

```
vector = df.iloc[0,1:5] # Access columns 1 to 5 of the 0-th row of the dataframe
print(vector)

vector = df.iloc[0,1:301] # Access columns 1 to 301 of the 0-th row of the dataframe
print(vector)
```

The output will look like:

```
1     0.073525
2    -0.031703
3     0.05401
4    -0.040015
Name: 0, dtype: object
1      0.073525
2     -0.031703
3       0.05401
4     -0.040015
5     -0.011894
         ...
296   -0.065746
297   -0.023745
298    0.012824
299    0.005162
300   -0.130008
Name: 0, Length: 300, dtype: object
```

Remember that similar to lists and array types in Python, indexing in pandas dataframes starts from 0. For example, the first row has an index equal to 0, the second an index equal to 1, and the n-th and index equal to n-1.

We now have stored the embedding for the first word in the dataset to the variable "vector".

6.1.5 Convert dataframe to numpy array

The majority of analysis and processing functions require numerical input to be of numpy array type. Let's convert the word embedding that we stored above in the variable "vector" to a numpy array of similar length.

```
import numpy as np

word_embedding = vector.to_numpy(dtype=float)

print(word_embedding)
```

The output will look like:

```
[ 0.073525 -0.031703 0.05401 -0.040015 -0.011894 0.002958 -0.065406
  0.079166 0.099932 -0.134534 0.016396 0.056965 -0.057572 0.010251
  0.050909 0.072398 -0.061399 0.162734 0.082844 -0.052345 0.019466
  0.058028 0.010818 0.030207 0.000889 -0.021392 0.030583 0.038453
 -0.165726 0.112191 -0.058096 -0.111807 -0.033361 0.094921 0.013087
 -0.010414 0.073115 -0.058642 0.099049 0.07455 0.035721 -0.059482
  0.06792 -0.058917 0.004033 0.040625 -0.106188 0.088392 0.009051
  0.019659 -0.026113 -0.09077 0.043568 -0.062696 -0.030408 0.03884
  0.004035 0.013125 0.02305 -0.041966 -0.025127 -0.034154 0.016506
 -0.079488 -0.039066 0.013653 -0.063375 0.091463 0.05381 0.024699
  0.062898 0.12055 -0.048261 -0.017775 0.029978 0.022461 0.094113
 -0.098662 0.011137 -0.080405 -0.056159 -0.055153 -0.0306 0.108031
  0.082347 0.058076 -0.113915 0.003602 0.019997 -0.033987 0.027162
  0.007531 -0.019095 -0.048414 -0.024189 0.04996 0.032027 -0.009882
```

```
  0.045991 0.055946 -0.028849 -0.019265 -0.041771 -0.06572 -0.024059
 -0.002533 0.02972  0.041398 0.123978 0.050968 0.072876 -0.091089
 -0.011735 -0.014177 0.058902 -0.145468 -0.12125 0.084168 -0.155788
 -0.009348 -0.015042 0.063526 0.039869 0.025571 0.044058 0.019486
  0.112362 0.031733 0.039299 -0.051543 -0.022537 -0.026686 -0.100046
  0.115309 0.008369 0.01551 -0.065277 0.031222 0.109851 -0.006308
 -0.016031 -0.038418 0.034439 0.025142 0.142227 0.04277 -0.01852
 -0.005247 -0.021285 -0.019829 0.131366 -0.002935 0.018499 0.040565
 -0.03535 -0.075773 -0.017759 0.033599 0.023961 -0.106251 -0.040328
 -0.012546 0.006421 -0.082573 -0.031654 -0.010218 0.053183 0.068255
 -0.027139 -0.062169 0.043021 0.027036 -0.006469 -0.142859 0.022744
  0.000512 -0.065334 -0.052299 -0.017929 0.03619 0.030412 0.022339
  0.080582 -0.007923 -0.006414 -0.024119 -0.039354 -0.00177 0.01856
  0.079291 -0.037962 0.004094 0.057353 -0.126054 0.039407 -0.047057
  0.028695 -0.041185 -0.042427 0.063292 -0.015259 -0.012919 0.029772
  0.001388 -0.046082 0.112506 -0.004109 0.020585 0.018128 -0.025253
  0.016204 0.035294 -0.042431 -0.014868 -0.141065 -0.073506 -0.021315
  0.067625 0.073685 0.023866 -0.010576 -0.042903 -0.071802 -0.071728
  0.019136 -0.087325 -0.042621 -0.064981 -0.013045 0.039378 -0.029022
 -0.054649 -0.008433 0.03112 -0.018196 -0.003567 0.021799 0.094146
 -0.017547 0.036818 0.012625 0.053266 -0.078154 -0.069845 0.019453
  0.047343 -0.033813 -0.001188 0.04178 0.010779 0.005534 0.010311
  0.093365 -0.010763 0.040849 0.021107 0.047443 -0.017278 -0.068958
  0.033909 -0.011073 0.067898 -0.00054 -0.012711 -0.042463 -0.015018
  0.051556 -0.023581 -0.064241 0.026324 -0.039002 -0.013781 -0.060861
 -0.048069 -0.02534 0.079009 0.102468 0.044573 -0.041997 0.081503
 -0.010224 -0.028924 0.008187 -0.062565 0.076221 -0.039867 0.029696
  0.094235 -0.072826 0.038318 0.048875 0.014754 -0.038729 0.033155
 -0.12039 -0.065746 -0.023745 0.012824 0.005162 -0.130008]
```

Note that when converting the pandas dataframe to a numpy array using the to_numpy() function, we indicated the type of numbers to be loaded. In this case, the numbers in the dataset are of float type. If for example the numbers were integers, we would have passed the "dtype=int" argument to the to_numpy() function.

6.2 Word embeddings

6.2.1 Load word embeddings

Let's now create a dictionary that will contain all the word embeddings available in the dataset in numpy array form. To test the dictionary, we will print the first 7 elements of the embeddings for the words "mother" and "boy".

```
word_embeddings = dict()
for i in range(count_row): # Iterate through all rows in dataframe (words)
    word = df.iloc[i,0] # Get word
    embedding = df.iloc[i,1:count_col].to_numpy(dtype=float) # Get embedding and convert to float
     numpy array
    word_embeddings[word] = embedding

print("mother ->",word_embeddings["mother"][0:7]) # Print the first 7 elements of the embedding for
     word "mother"
print("boy   ->",word_embeddings["boy"][0:7]) # Print the first 7 elements of the embedding for word
     "boy"
```

The output will look like:

```
mother -> [ 0.007815 0.026617 -0.036383 -0.051246 0.000183 0.071259 0.017416]
boy    -> [ 0.040253 0.001048 0.023576 0.003103 -0.027837 0.035486 0.04829 ]
```

6.2.2 Distance of word embeddings

Let's now compute and print a matrix with the pairwise cosine distances of the embeddings of the words "mother", "boy", "sister", "family", "home", "rabbit", and "eyes".

```python
from scipy.spatial import distance

test_words = ["mother","boy","sister","family","home","rabbit","eyes"]

print("%6s" % "", end="")
for word in test_words:
    print("\t%6s" % word, end="")
print("")
for word1 in test_words:
    print("%6s" % word1, end="")
    for word2 in test_words:
        print("\t%1.4f" % distance.cosine(word_embeddings[word1],word_embeddings[word2]), end="")
    print("")
```

The output will look like:

```
        mother     boy  sister  family    home  rabbit    eyes
mother  0.0000  0.4318  0.3882  0.4893  0.7175  0.7303  0.7059
   boy  0.4318  0.0000  0.6213  0.7388  0.8153  0.6080  0.6594
sister  0.3882  0.6213  0.0000  0.6224  0.7767  0.8759  0.8403
family  0.4893  0.7388  0.6224  0.0000  0.7678  0.8327  0.8262
  home  0.7175  0.8153  0.7767  0.7678  0.0000  0.8341  0.8851
rabbit  0.7303  0.6080  0.8759  0.8327  0.8341  0.0000  0.6599
  eyes  0.7059  0.6594  0.8403  0.8262  0.8851  0.6599  0.0000
```

Notice that the words "mother", "sister", "boy" and "family" have smaller distances among each other, compared to the other words. Considering that they have contextual relations, it is an indication that these word embeddings are able to encode contextual information about the words.

6.3 Document embeddings

Let's now use the available word embeddings in order to create the document embeddings for the following documents:

1. My mother was sitting on the bed

2. The night looked remarkable at the beginning

There are various ways to create document embeddings. One of the most common approaches is to compute the embeddings for each word in a document and then compute the average embedding (applied element-wise) across the embeddings of its constituent words.

We will first tokenise the two documents:

```python
from nltk import word_tokenize # Import the word_tokenize function from NLTK

text1 = "My mother was sitting on the bed"
text2 = "The night looked remarkable at the beginning"

tokens1 = word_tokenize(text1.lower()) # Tokenise "text1" into words
tokens2 = word_tokenize(text2.lower()) # Tokenise "text1" into words

words_list1 = []
for word in tokens1:
    words_list1.append(word)

print(text1,"->",words_list1)

words_list2 = []
for word in tokens2:
    words_list2.append(word)

print(text2,"->",words_list2)
```

The output will look like:

```
My mother was sitting on the bed -> ['my', 'mother', 'was', 'sitting', 'on', 'the', 'bed']
The night looked remarkable at the beginning -> ['the', 'night', 'looked', 'remarkable', 'at', 'the',
    'beginning']
```

We will now use the dictionary of the word embeddings that we created before in order to retrieve the embedding for each word in the two documents.

```python
print("Text1 word embeddings:")
for word in words_list1:
    try:
        print(word,"->",word_embeddings[word][0:4]) # Print first 4 elements of embedding
    except:
        print(word,"-> n/a")

print("\nText2 word embeddings:")
for word in words_list2:
    try:
        print(word,"->",word_embeddings[word][0:4]) # Print first 4 elements of embedding
    except:
        print(word,"-> n/a")
```

The output will look like:

```
Text1 word embeddings:
my -> n/a
mother -> [ 0.007815 0.026617 -0.036383 -0.051246]
was -> n/a
sitting -> [9.7000e-05 4.1846e-02 8.3230e-02 5.5060e-03]
on -> [ 0.005252 -0.002234 -0.0648 -0.001852]
the -> [ 0.016258 -0.013271 -0.007168 -0.083179]
bed -> [ 0.0882   0.056767 -0.021443 0.014364]

Text2 word embeddings:
the -> [ 0.016258 -0.013271 -0.007168 -0.083179]
night -> [ 0.06248 -0.051441 -0.023803 0.038181]
looked -> [-0.011535 0.016638 0.063261 -0.033455]
remarkable -> [ 0.065266 -0.021301 0.060876 0.068074]
at -> n/a
the -> [ 0.016258 -0.013271 -0.007168 -0.083179]
beginning -> [-0.017209 0.131408 0.024707 -0.036714]
```

As you can see above, some of the words from the two documents are missing from the word embeddings dataset that we have. In this case, we will ignore the words with the missing embeddings and process each document as if these words do not exist.

6.3.1   Compute document embedding (Mean word embedding)

Let's now create a function that given a document in the form of a list of words, a dictionary of word embeddings and the size k of the requested document embedding, will return a document embedding of size k, computed as the first k elements of the mean word embedding of the document's constituent words. Then, we will print the embedding of the two documents for $k = 5$ and $k = 20$.

```python
def get_document_embedding(word_list,word_embeddings,k):
    document_embedding = np.zeros(k,dtype=float) # Create embedding of k zero-valued elements
    valid_words = 0
    for word in word_list:
        try:
            document_embedding = document_embedding + word_embeddings[word][0:k] # Add word embedding
     to partial sum
            valid_words += 1
        except:
            pass # If word embedding is not available, then ignore the word
    document_embedding = document_embedding / valid_words # Divide all elements by number of valid
     words to get the mean
```

```
    return document_embedding

print("Text1 embedding (k=5) ->",get_document_embedding(words_list1,word_embeddings,5),"\n")
print("Text2 embedding (k=5) ->",get_document_embedding(words_list2,word_embeddings,5),"\n")

print("Text1 embedding (k=20) ->",get_document_embedding(words_list1,word_embeddings,20),"\n")
print("Text2 embedding (k=20) ->",get_document_embedding(words_list2,word_embeddings,20),"\n")
```

The output will look like:

```
Text1 embedding (k=5) -> [ 0.0235244 0.021945 -0.0093128 -0.0232814 0.0077766]

Text2 embedding (k=5) -> [ 0.02191967 0.008127 0.01845083 -0.021712 -0.008242 ]

Text1 embedding (k=20) -> [ 0.0235244 0.021945 -0.0093128 -0.0232814 0.0077766 0.0664142
 -0.0294452 0.0385194 0.0415822 0.0027638 0.038138 -0.0464678
  0.0299128 -0.0162724 0.0267094 -0.0193898 -0.0579416 0.0045326
  0.0272054 0.002882 ]

Text2 embedding (k=20) -> [ 0.02191967 0.008127 0.01845083 -0.021712 -0.008242 0.05405583
 -0.00840717 0.07847117 0.017932 -0.01513817 0.035567 -0.04270783
  0.05314033 -0.03148433 -0.0126355 -0.016384 -0.08776767 0.01809233
  0.06723833 -0.00525483]
```

6.3.2   Cosine distance of document embeddings

Let's now compute the cosine distance between the document embeddings of the two documents for $k = 5, 20, 15, 300$.

```
print("Cosine distance of text1 and text2 for k =
    5:",distance.cosine(get_document_embedding(words_list1,word_embeddings,5),
    get_document_embedding(words_list2,word_embeddings,5)))
print("Cosine distance of text1 and text2 for k =
    20:",distance.cosine(get_document_embedding(words_list1,word_embeddings,20),
    get_document_embedding(words_list2,word_embeddings,20)))
print("Cosine distance of text1 and text2 for k =
    150:",distance.cosine(get_document_embedding(words_list1,word_embeddings,150),
    get_document_embedding(words_list2,word_embeddings,150)))
print("Cosine distance of text1 and text2 for k =
    300:",distance.cosine(get_document_embedding(words_list1,word_embeddings,300),g
    et_document_embedding(words_list2,word_embeddings,300)))
```

The output will look like:

```
Cosine distance of text1 and text2 for k = 5: 0.3855626793684508
Cosine distance of text1 and text2 for k = 20: 0.16427167205198012
Cosine distance of text1 and text2 for k = 150: 0.28613327209110107
Cosine distance of text1 and text2 for k = 300: 0.3327962278223764
```

6.4   Classification using document embeddings

Let's now use document embeddings based on our word embeddings dataset to classify movie reviews as having positive or negative sentiment. We will first load the movie reviews dataset, which consists of 20 files, 10 files named pos_XX.txt that each contains one movie review with positive sentiment, and 10 files named neg_XX.txt that each contains one movie review with negative sentiment. XX is an identification number from 01 to 10.

```
text = []
label = []

for i in range(1,11):
    filename_pos = "pos_%02d.txt" % i # Create string with the filename for positive sentiment reviews
    filename_neg = "neg_%02d.txt" % i # Create string with the filename for negative sentiment reviews
```

```python
    print(filename_pos)
    print(filename_neg)
    # Open positive sentiment file
    f = open(filename_pos, "r") # Opens the file for reading only ("r")
    text.append(f.read())
    f.close() # Close the file
    label.append("pos") # Add positive sentiment label to labels list
    # Open negative sentiment file
    f = open(filename_neg, "r") # Opens the file for reading only ("r")
    text.append(f.read())
    f.close() # Close the file
    label.append("neg") # Add negative sentiment label to labels list

print("No of texts:",len(text))
print("No of labels:",len(label))
```

The output will look like:

```
pos_01.txt
neg_01.txt
pos_02.txt
neg_02.txt
pos_03.txt
neg_03.txt
pos_04.txt
neg_04.txt
pos_05.txt
neg_05.txt
pos_06.txt
neg_06.txt
pos_07.txt
neg_07.txt
pos_08.txt
neg_08.txt
pos_09.txt
neg_09.txt
pos_10.txt
neg_10.txt
No of texts: 20
No of labels: 20
```

Now let's compute the document embedding for each document, using the function that we created before.

```python
from nltk import word_tokenize
from string import punctuation

punctuation_list = list(punctuation)

text_embeddings = []
for i in range(len(text)): # Iterate through all texts
    tokens = word_tokenize(text[i].lower()) # Tokenise "text" into words
    words_list = []
    for word in tokens:
        if(word not in punctuation_list):
            words_list.append(word)
    text_embeddings.append(get_document_embedding(words_list,word_embeddings,300))

for i in range(len(text)): # Iterate through all texts
    print(i,text_embeddings[i][0:5],"->",label[i]) # Print the first 5 elements of each document
     embedding
```

The output will look like:

```
0 [-0.00144049 -0.01854695 0.01171497 -0.05752111 -0.03104189] -> pos
1 [-0.00351737 -0.02884196 0.00754981 -0.03615796 -0.028105 ] -> neg
```

```
2 [ 0.01211377 -0.01454308 0.01602908 -0.06364362 -0.02417946] -> pos
3 [ 0.0013629 -0.02655979 0.00252164 -0.05645312 -0.02949102] -> neg
4 [-0.025445 -0.03421304 0.012184 -0.04374732 -0.0253342 ] -> pos
5 [-0.0044761 -0.0279336 0.0044961 -0.04975507 -0.0246164 ] -> neg
6 [ 0.01035458 -0.02045229 0.00215542 -0.0624085 -0.04229192] -> pos
7 [-0.00919744 -0.02292037 0.01568922 -0.0474781 -0.02196531] -> neg
8 [-0.01577481 -0.02448513 0.00704729 -0.04966781 -0.02111129] -> pos
9 [ 0.00515638 -0.01661385 0.00617285 -0.04173085 -0.04571008] -> neg
10 [ 0.00058903 -0.01620987 0.0142112 -0.0417965 -0.0315724 ] -> pos
11 [ 0.0067564 -0.01518371 0.00774966 -0.04712666 -0.03590557] -> neg
12 [-0.00779024 -0.02070952 0.00332203 -0.05560338 -0.0289061 ] -> pos
13 [-0.00285736 -0.0158494 0.008223 -0.05303438 -0.02770218] -> neg
14 [-0.011991 -0.02553873 -0.01264482 -0.05226082 -0.02133141] -> pos
15 [-0.01091673 -0.02792297 0.00905532 -0.0435513 -0.03137696] -> neg
16 [-0.00434465 -0.01706052 -0.00136974 -0.04984022 -0.03162339] -> pos
17 [-0.01131559 -0.01907496 -0.00221819 -0.05785556 -0.01828148] -> neg
18 [ 0.00109655 -0.02176745 0.00632607 -0.0452914 -0.03907748] -> pos
19 [ 0.00515789 -0.01992803 0.02873918 -0.04443174 -0.01608311] -> neg
```

Note that the sentiment labels are of type string. Many machine learning packages and functions require the labels to be in numerical format. To achieve this, we will encode our sentiment labels to numbers.

```python
from sklearn import preprocessing

le = preprocessing.LabelEncoder() # Create labelEncoder

labels_encoded=le.fit_transform(label) # Encode labels to numbers
print(label,"->",labels_encoded)
```

The output will look like:

```
['pos', 'neg', 'pos', 'neg', 'pos', 'neg', 'pos', 'neg', 'pos', 'neg', 'pos', 'neg', 'pos', 'neg',
    'pos', 'neg', 'pos', 'neg', 'pos', 'neg'] -> [1 0 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1 0 1 0]
```

As you can see above, the negative ("neg") label is now denoted by 0 and the positive ("pos") label by 1.

6.4.1 Classification with the k Nearest Neighbour algorithm (kNN)

Let's use the k Nearest Neighbours algorithm for 3 nearest neighbours in order to classify the movie reviews in positive or negative. We will first divide our dataset into a training set consisting of 14 samples (7 negative and 7 positive) and a test set consisting of 6 samples (3 negative and 3 positive). We will then train the kNN model on the training set and compute the classification performance on the test set.

```python
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix
from sklearn.metrics import accuracy_score

# Divide dataset to training (14 samples - 7 positive,7 negative) and test (6 samples - 3 positive, 3
    negative)
training_features = text_embeddings[0:14]
training_labels = labels_encoded[0:14]
test_features = text_embeddings[14:20]
test_labels = labels_encoded[14:20]


model = KNeighborsClassifier(n_neighbors=3)

# Train the model using the training set
model.fit(training_features,training_labels)

#Predict Output
predicted= model.predict(test_features)
print("Prediction :",predicted)
print("True labels:",test_labels)
```

```
cm = confusion_matrix(test_labels, predicted) # Create confusion matrix
accuracy = accuracy_score(test_labels, predicted) # Compute classification accuracy

print("Confusion matrix:\n%s" % cm)
print("Accuracy: %.2f%s" % (accuracy*100,"%"))
```

The output will look like:

```
Prediction : [1 0 1 1 1 1]
True labels: [1 0 1 0 1 0]
Confusion matrix:
[[1 2]
 [0 3]]
Accuracy: 66.67%
```

As you can see above, the kNN classifier for 3 nearest neighbours was able to classify correctly 4 out of the 6 samples in our test set, reaching a classification accuracy of 66.67%.

6.4.2 Classification with Linear Support Vector Machines (SVM)

Let's repeat our experiment using a Linear Support Vector Machine (SVM) model.

```
#Import svm model
from sklearn import svm

#Create a svm Classifier
clf = svm.SVC(kernel='linear') # Linear Kernel

#Train the model using the training set
clf.fit(training_features,training_labels)

#Predict the response for test dataset
predicted = clf.predict(test_features)
print("Prediction :",predicted)
print("True labels:",test_labels)

cm = confusion_matrix(test_labels, predicted)
accuracy = accuracy_score(test_labels, predicted)

print("Confusion matrix:\n%s" % cm)
print("Accuracy: %.2f%s" % (accuracy*100,"%"))
```

The output will look like:

```
Prediction : [1 0 1 1 0 0]
True labels: [1 0 1 0 1 0]
Confusion matrix:
[[2 1]
 [1 2]]
Accuracy: 66.67%
```

As you can see above, the SVM classifier achieved the same classification accuracy as the kNN classifier. However, note that it miss-classified one sample from each class, whereas the kNN classifier miss-classified two samples from the same class. Remember that metrics such as the classification accuracy are not sufficient to provide a complete overview of a machine learning model's performance.

6.4.3 Classification using Feed-Forward Neural Networks

Let's now use a Feed-Forward Neural Network for the same task. We will define a neural network that has an input layer of size 300 (300 being the size of the embedding), one hidden layer with 5 neurons and an output layer of 2 neurons (2 being the number of classes). A ReLU activation function will be used for the hidden layer, while a softmax activation function will be used for the output layer in order to convert the network's output to class probabilities.

```
import tensorflow as tf
```

```
EMBEDDING_SIZE = 300

model = tf.keras.models.Sequential() # Create a sequential neural network model
model.add(tf.keras.Input(shape=(EMBEDDING_SIZE,))) # The input size will be equal to the embedding
    size
model.add(tf.keras.layers.Dense(5, activation='relu')) # 1st hidden layer has a size of 5 neurons
model.add(tf.keras.layers.Dense(2, activation='softmax')) # Output layer has a size of 2 neurons

model.summary() # Print a summary of the model
```

The output will look like:

```
Model: "sequential"

----------------------------------------------------------------
 Layer (type)              Output Shape            Param #
================================================================
 dense (Dense)             (None, 5)               1505


 dense_1 (Dense)           (None, 2)               12


================================================================
Total params: 1,517
Trainable params: 1,517
Non-trainable params: 0

----------------------------------------------------------------
```

Now let's train the network for 2000 epochs, using the Adam optimiser, a learning rate of $\eta = 0.001$, and cross entropy loss as the loss function.

```python
import matplotlib.pyplot as plt

# Define Adam optimiser and set learning rate to 0.001
opt = tf.keras.optimizers.Adam(learning_rate=0.001)

# Compile model using sparse categorical cross entropy loss as loss function and the Adam optimiser
model.compile(optimizer=opt, loss='sparse_categorical_crossentropy', metrics='accuracy')

EPOCHS = 2000 # Train model for 2000 epochs

# Convert data to numpy arrays
X_train = np.array(training_features)
y_train = np.array(training_labels)
X_test = np.array(test_features)
y_test = np.array(test_labels)

# Train model and store the loss and the metrics for each epoch in the history variable
history = model.fit(x=X_train, y=y_train, validation_data=(X_test,y_test),epochs=EPOCHS,verbose=0)

# Plot training progress
plt.style.use('ggplot')
fig, (ax1, ax2) = plt.subplots(2, figsize=(12, 6), sharex=True)

ax1.set_ylim([0, 1.01])
ax1.plot(history.history['val_accuracy'],'b')
ax1.set_ylabel("Accuracy")
ax1.plot(history.history['accuracy'],'r')
ax1.legend(['Validation','Training'])

ax2.set_ylim([0, 1])
ax2.plot(history.history['val_loss'],'b')
ax2.plot(history.history['loss'],'r')
ax2.set_ylabel("Loss")
ax2.set_xlabel("Epochs")
ax2.legend(['Validation','Training'])
```

```python
print("Max validation accuracy: %.2f%s" % (max(history.history['val_accuracy'])*100,"%"))

for i in range(EPOCHS):
    if(history.history['val_accuracy'][i] == max(history.history['val_accuracy'])):
        print("Max accuracy reached at epoch",i)
        break
```

The output will look like:

```
Max validation accuracy: 83.33%
Max accuracy reached at epoch 410
```


Note that the maximum accuracy and the epoch that it was reached may differ a bit every time your rerun the code due to the random initialisation of the neural network's weights.

6.5   Exercises

Exercise 6.1 Using the word embeddings from testvectors.csv, compute the word embeddings of size k for each word in the text "Once upon a time, my family was living in an ancient castle.". Then compute the pairwise cosine distance between the words in the text and visualise them as a heatmap. Repeat this for $k = 5, 50, 300$. Ignore words for which an embedding is not available.

Exercise 6.2 Repeat the classification task from Section 6.4.1 for an embedding of size $k = 50$ and an embedding of size $k = 200$.

Exercise 6.3 Improve the feed forward neural network's architecture from Section 6.4.3 in order to make it reach the maximum accuracy earlier (in less epochs). Note: Try to experiment with adding more layers, changing the size of the layers, using different activation functions for the hidden layers, training for more epochs, changing the learning rate, etc.

Appendix: Test files

A.1   movies.xml

```xml
<movies>
  <movie id="1">
    <title>And Now for Something Completely Different</title>
    <released>1971</released>
  </movie>
  <movie id="2">
    <title>Monty Python and the Holy Grail</title>
    <released>1974</released>
  </movie>
  <movie id="3">
    <title>Monty Python's Life of Brian</title>
    <released>1979</released>
  </movie>
  <movie id="4">
    <title>Monty Python Live at the Hollywood Bowl</title>
    <released>1982</released>
  </movie>
  <movie id="5">
    <title>Monty Python's The Meaning of Life</title>
    <released>1983</released>
  </movie>
</movies>
```

A.2   links.html

```html
<!doctype html>
<html lang="en-GB">

<head>
    <meta name="viewport" content="width=device-width, initial-scale=1">
    <title>Test page for Text Mining and Language Analytics</title>
</head>

<body>
    <h1>Test page for Text Mining and Language Analytics</h1>
    <a href="https://www.durham.ac.uk/departments/academic/computer-science/">Link 1</a><br/>
    <a href="https://www.gov.uk/government/organisations/hm-revenue-customs">Link 2</a><br/>
    <a href="https://www.gov.uk/">Link 3</a><br/>
    <a href="https://www.dur.ac.uk/">Link 4</a><br/>
    <a href="https://www.nhs.uk/">Link 5</a><br/>
</body>

</html>
```

A.3   emails.txt

john+acme.co@hotmail.com
bob@gmail.com
tom@durham.ac.uk
jerry@durham.ac.uk
scrooge@durham.ac.uk
donald@yahoo.co.uk
huey@yahoo.co.uk
dewey@gmail.com
louie.duck@durham.ac.uk
gyro.gearloose@yahoo.co.uk
bart@yahoo.co.uk
homer@gmail.com
stan@hotmail.com
kyle-broflovski@durham.ac.uk
eric@yahoo.co.uk
kenny@gmail.com
butters@durham.ac.uk
wendy@hotmail.com
randy_marsh@durham.ac.uk
chef@gmail.com

A.4   cds.xml

```xml
<?xml version="1.0" encoding="ISO8859-1" ?>
<CATALOG>
<CD>
<TITLE>Empire Burlesque</TITLE>
<ARTIST>Bob Dylan</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>Columbia</COMPANY>
<PRICE>10.90</PRICE>
<YEAR>1985</YEAR>
</CD>
<CD>
<TITLE>Hide your heart</TITLE>
<ARTIST>Bonnie Tylor</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>CBS Records</COMPANY>
<PRICE>9.90</PRICE>
<YEAR>1988</YEAR>
</CD>
<CD>
<TITLE>Greatest Hits</TITLE>
<ARTIST>Dolly Parton</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>RCA</COMPANY>
<PRICE>9.90</PRICE>
<YEAR>1982</YEAR>
</CD>
<CD>
<TITLE>Still got the blues</TITLE>
<ARTIST>Gary More</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Virgin redords</COMPANY>
<PRICE>10.20</PRICE>
<YEAR>1990</YEAR>
</CD>
<CD>
<TITLE>Eros</TITLE>
<ARTIST>Eros Ramazzotti</ARTIST>
```

```
<COUNTRY>EU</COUNTRY>
<COMPANY>BMG</COMPANY>
<PRICE>9.90</PRICE>
<YEAR>1997</YEAR>
</CD>
<CD>
<TITLE>One night only</TITLE>
<ARTIST>Bee Gees</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Polydor</COMPANY>
<PRICE>10.90</PRICE>
<YEAR>1998</YEAR>
</CD>
<CD>
<TITLE>Sylvias Mother</TITLE>
<ARTIST>Dr.Hook</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>CBS</COMPANY>
<PRICE>8.10</PRICE>
<YEAR>1973</YEAR>
</CD>
<CD>
<TITLE>Maggie May</TITLE>
<ARTIST>Rod Stewart</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Pickwick</COMPANY>
<PRICE>8.50</PRICE>
<YEAR>1990</YEAR>
</CD>
<CD>
<TITLE>Romanza</TITLE>
<ARTIST>Andrea Bocelli</ARTIST>
<COUNTRY>EU</COUNTRY>
<COMPANY>Polydor</COMPANY>
<PRICE>10.80</PRICE>
<YEAR>1996</YEAR>
</CD>
<CD>
<TITLE>When a man loves a woman</TITLE>
<ARTIST>Percy Sledge</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>Atlantic</COMPANY>
<PRICE>8.70</PRICE>
<YEAR>1987</YEAR>
</CD>
<CD>
<TITLE>Black angel</TITLE>
<ARTIST>Savage Rose</ARTIST>
<COUNTRY>EU</COUNTRY>
<COMPANY>Mega</COMPANY>
<PRICE>10.90</PRICE>
<YEAR>1995</YEAR>
</CD>
<CD>
<TITLE>1999 Grammy Nominees</TITLE>
<ARTIST>Many</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>Grammy</COMPANY>
<PRICE>10.20</PRICE>
<YEAR>1999</YEAR>
</CD>
<CD>
<TITLE>For the good times</TITLE>
<ARTIST>Kenny Rogers</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Mucik Master</COMPANY>
```

```
<PRICE>8.70</PRICE>
<YEAR>1995</YEAR>
</CD>
<CD>
<TITLE>Big Willie style</TITLE>
<ARTIST>Will Smith</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>Columbia</COMPANY>
<PRICE>9.90</PRICE>
<YEAR>1997</YEAR>
</CD>
<CD>
<TITLE>Tupelo Honey</TITLE>
<ARTIST>Van Morrison</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Polydor</COMPANY>
<PRICE>8.20</PRICE>
<YEAR>1971</YEAR>
</CD>
<CD>
<TITLE>Soulsville</TITLE>
<ARTIST>Jorn Hoel</ARTIST>
<COUNTRY>Norway</COUNTRY>
<COMPANY>WEA</COMPANY>
<PRICE>7.90</PRICE>
<YEAR>1996</YEAR>
</CD>
<CD>
<TITLE>The very best of</TITLE>
<ARTIST>Cat Stevens</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Island</COMPANY>
<PRICE>8.90</PRICE>
<YEAR>1990</YEAR>
</CD>
<CD>
<TITLE>Stop</TITLE>
<ARTIST>Sam Brown</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>A and M</COMPANY>
<PRICE>8.90</PRICE>
<YEAR>1988</YEAR>
</CD>
<CD>
<TITLE>Bridge of Spies</TITLE>
<ARTIST>T`Pau</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Siren</COMPANY>
<PRICE>7.90</PRICE>
<YEAR>1987</YEAR>
</CD>
<CD>
<TITLE>Private Dancer</TITLE>
<ARTIST>Tina Turner</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>Capitol</COMPANY>
<PRICE>8.90</PRICE>
<YEAR>1983</YEAR>
</CD>
<CD>
<TITLE>Midt om natten</TITLE>
<ARTIST>Kim Larsen</ARTIST>
<COUNTRY>EU</COUNTRY>
<COMPANY>Medley</COMPANY>
<PRICE>7.80</PRICE>
<YEAR>1983</YEAR>
```

```
</CD>
<CD>
<TITLE>Pavarotti Gala Concert</TITLE>
<ARTIST>Luciano Pavarotti</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>DECCA</COMPANY>
<PRICE>9.90</PRICE>
<YEAR>1991</YEAR>
</CD>
<CD>
<TITLE>The dock of the bay</TITLE>
<ARTIST>Otis Redding</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>Atlantic</COMPANY>
<PRICE>7.90</PRICE>
<YEAR>1987</YEAR>
</CD>
<CD>
<TITLE>Picture book</TITLE>
<ARTIST>Simply Red</ARTIST>
<COUNTRY>EU</COUNTRY>
<COMPANY>Elektra</COMPANY>
<PRICE>7.20</PRICE>
<YEAR>1985</YEAR>
</CD>
<CD>
<TITLE>Red</TITLE>
<ARTIST>The Communards</ARTIST>
<COUNTRY>UK</COUNTRY>
<COMPANY>London</COMPANY>
<PRICE>7.80</PRICE>
<YEAR>1987</YEAR>
</CD>
<CD>
<TITLE>Unchain my heart</TITLE>
<ARTIST>Joe Cocker</ARTIST>
<COUNTRY>USA</COUNTRY>
<COMPANY>EMI</COMPANY>
<PRICE>8.20</PRICE>
<YEAR>1987</YEAR>
</CD>
</CATALOG>
```

A.5   alice.txt

Alice was beginning to get very tired of sitting by her sister on the bank, and of having nothing to do: once or twice she had peeped into the book her sister was reading, but it had no pictures or conversations in it, "and what is the use of a book," thought Alice "without pictures or conversations?"

So she was considering in her own mind (as well as she could, for the hot day made her feel very sleepy and stupid), whether the pleasure of making a daisy-chain would be worth the trouble of getting up and picking the daisies, when suddenly a White Rabbit with pink eyes ran close by her.

There was nothing so very remarkable in that; nor did Alice think it so very much out of the way to hear the Rabbit say to itself, "Oh dear! Oh dear! I shall be late!" (when she thought it over afterwards, it occurred to her that she ought to have wondered at this, but at the time it all seemed quite natural); but when the Rabbit actually took a watch out of its waistcoat-pocket, and looked at it, and then hurried on, Alice started to her feet, for it flashed across her mind that she had never before seen a rabbit with either a waistcoat-pocket, or a watch to take out of it, and burning with curiosity, she ran across the field after it, and fortunately was just in time to see it pop down a large rabbit-hole under the hedge.

A.6   dune.txt

In the week before their departure to Arrakis, when all the final scurrying about had reached a nearly unbearable frenzy, an old crone came to visit the mother of the boy, Paul.

It was a warm night at Castle Caladan, and the ancient pile of stone that had served the Atreides family as home for twenty-six generations bore that cooled-sweat feeling it acquired before a change in the weather.

The old woman was let in by the side door down the vaulted passage by Paul's room and she was allowed a moment to peer in at him where he lay in his bed.

By the half-light of a suspensor lamp, dimmed and hanging near the floor, the awakened boy could see a bulky female shape at his door, standing one step ahead of his mother. The old woman was a witch shadow - hair like matted spiderwebs, hooded 'round darkness of features, eyes like glittering jewels.